



10-708 Probabilistic Graphical Models

Machine Learning Department
School of Computer Science
Carnegie Mellon University



MAP Inference by MILP

Matt Gormley
Lecture 9
Mar. 1, 2021

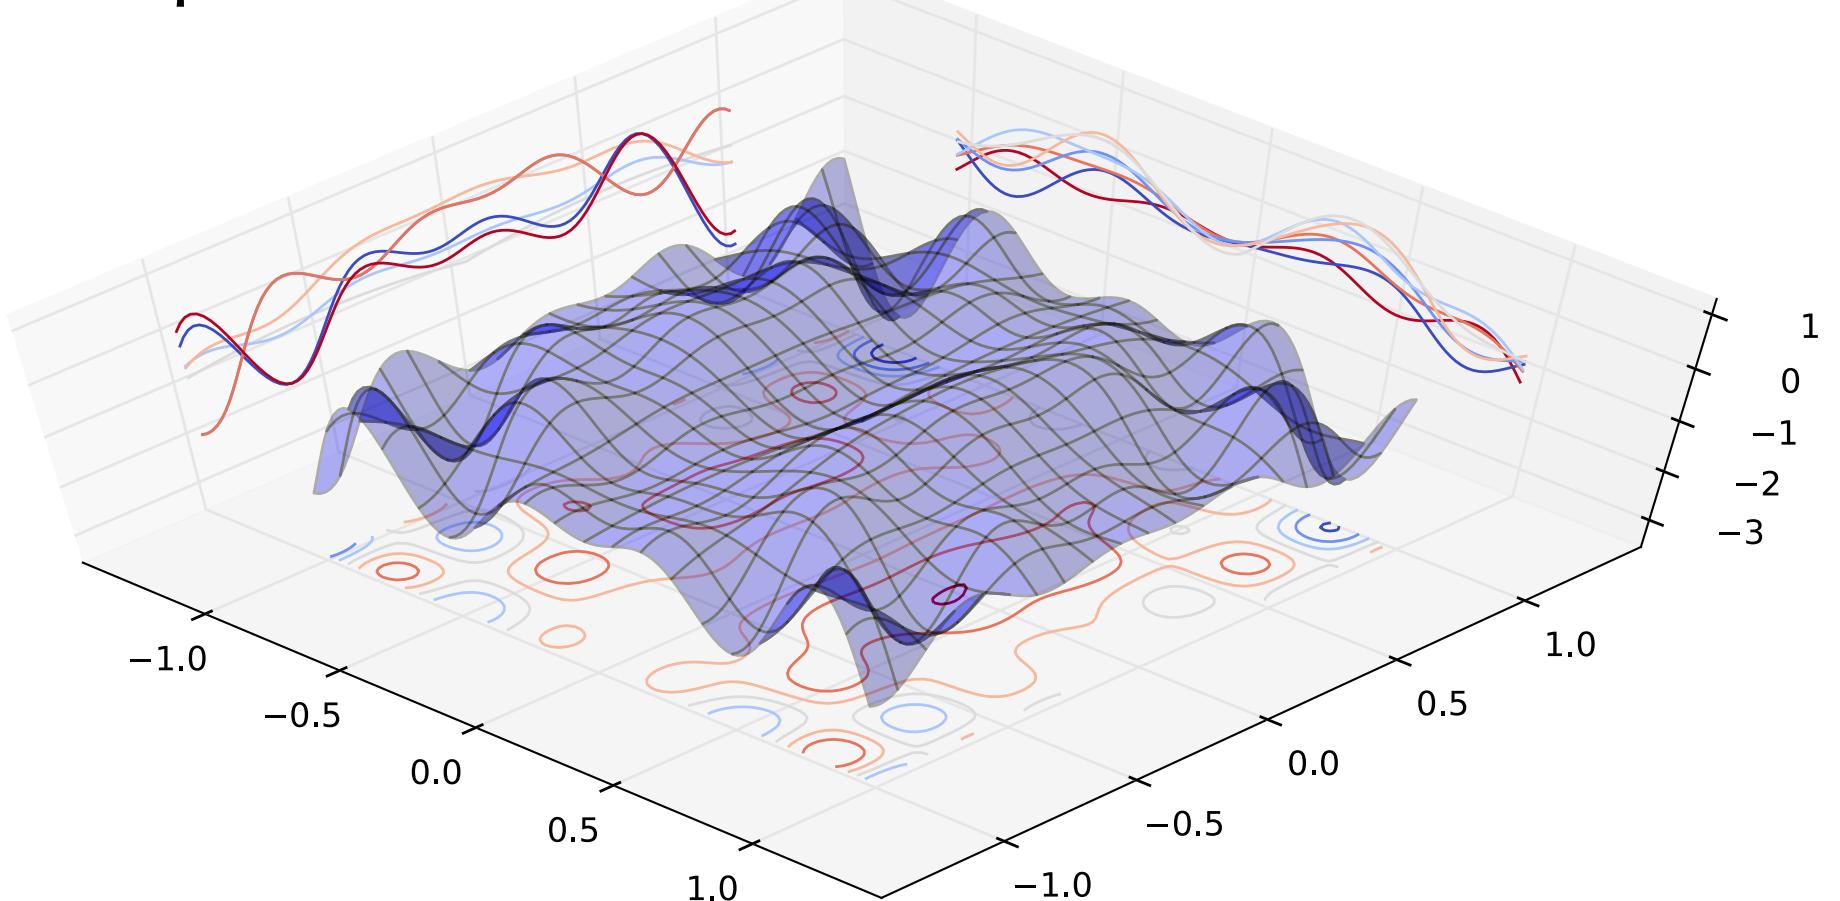
Reminders

- **Quiz 1:**
 - Fri, Mar. 05 during class time
 - Material: Lectures 1 – 8 only
 - Logistics: see forthcoming Piazza post
- **Homework 2: Exact inference and supervised learning (CRF+RNN)**
 - Out: Wed, Feb. 24
 - Due: Wed, Mar. 10 at 11:59pm

BRANCH AND BOUND

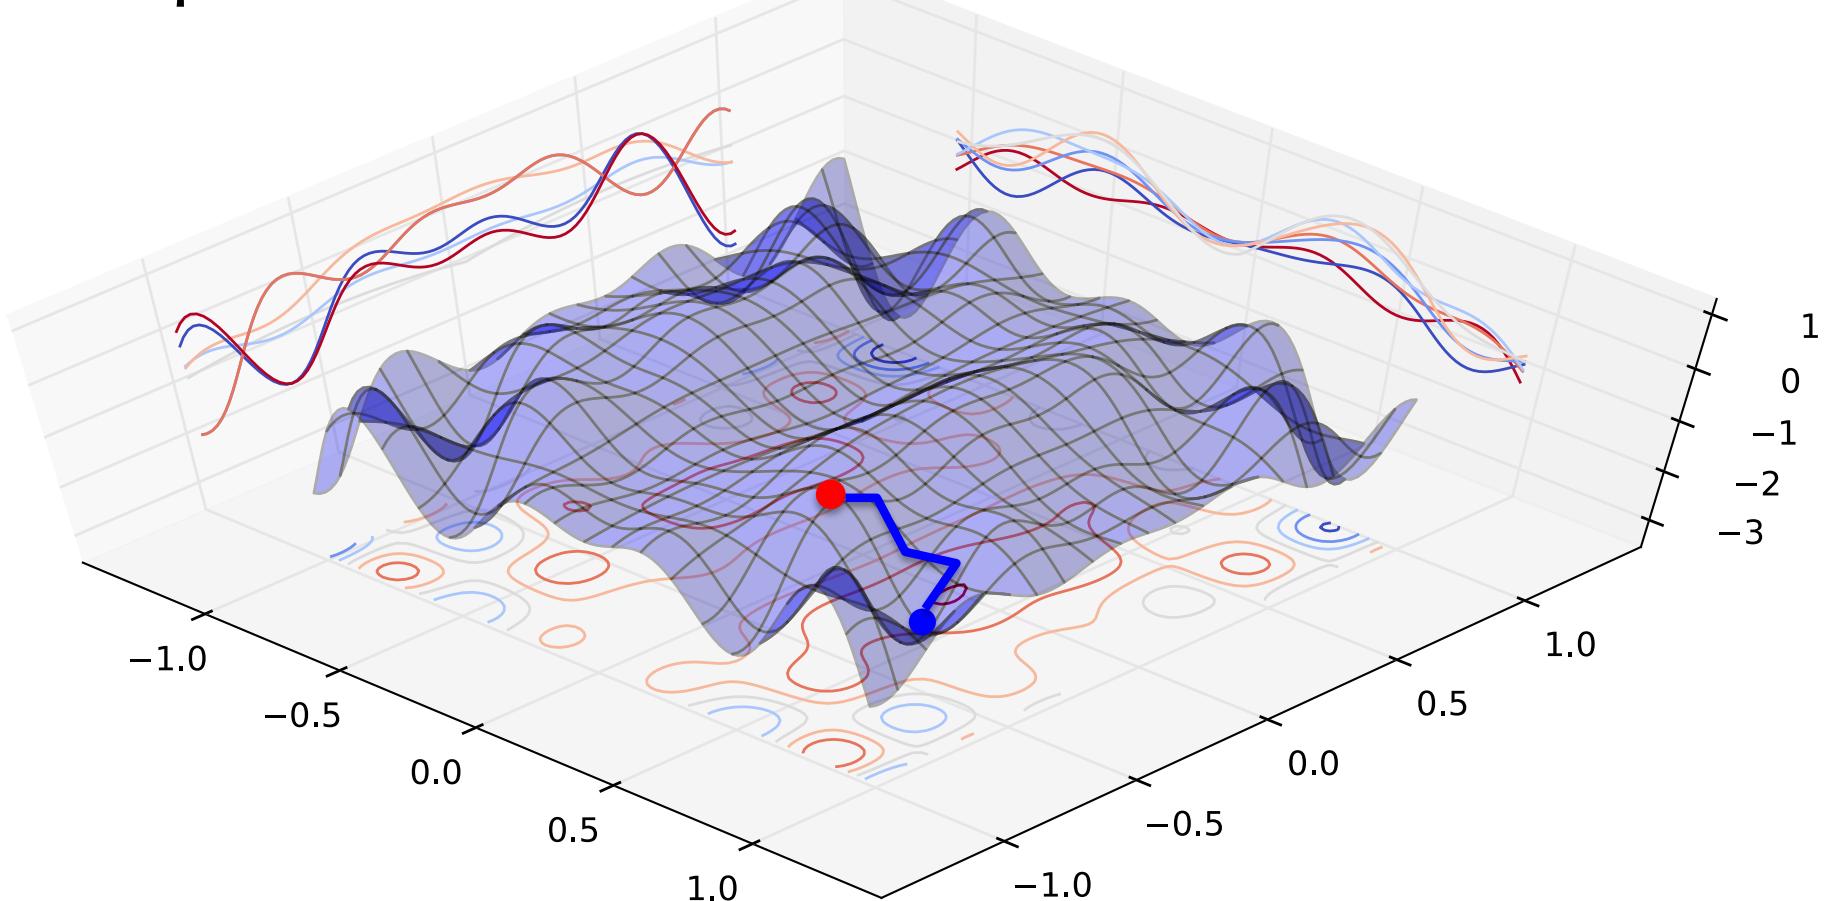
Background: Nonconvex Global Optimization

Goal: optimize over the blue surface.



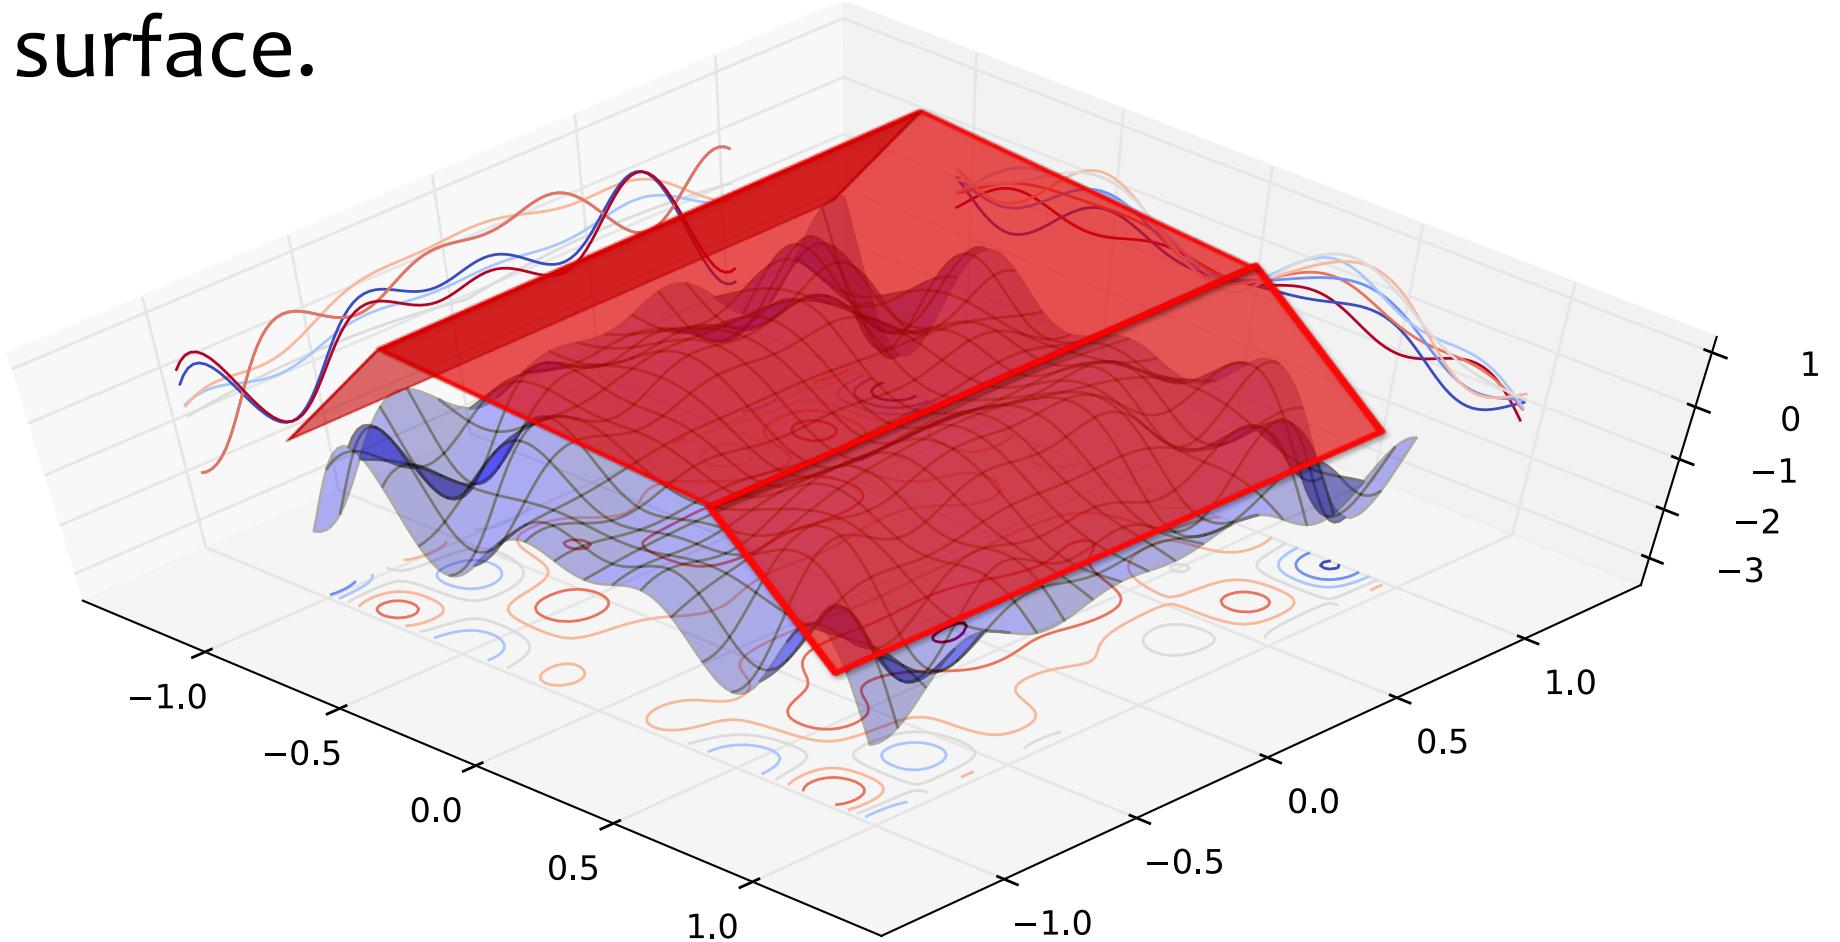
Background: Nonconvex Global Optimization

Goal: optimize over the blue surface.



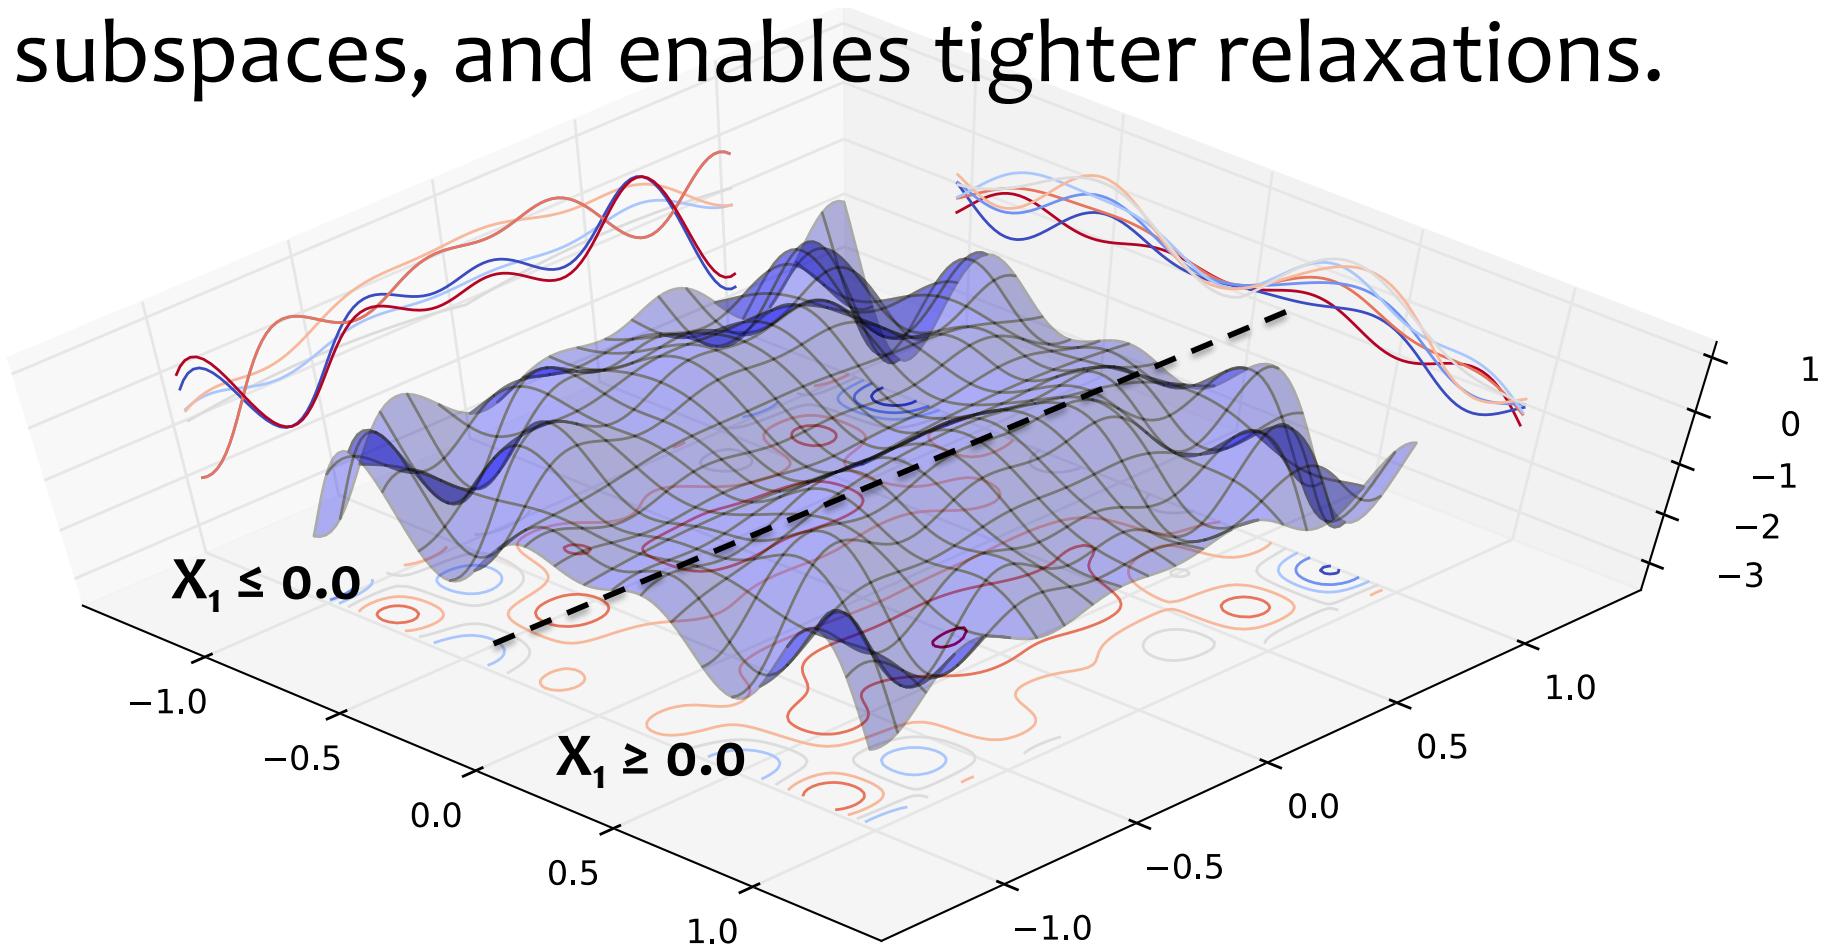
Background: Nonconvex Global Optimization

Relaxation: provides an upper bound on the surface.



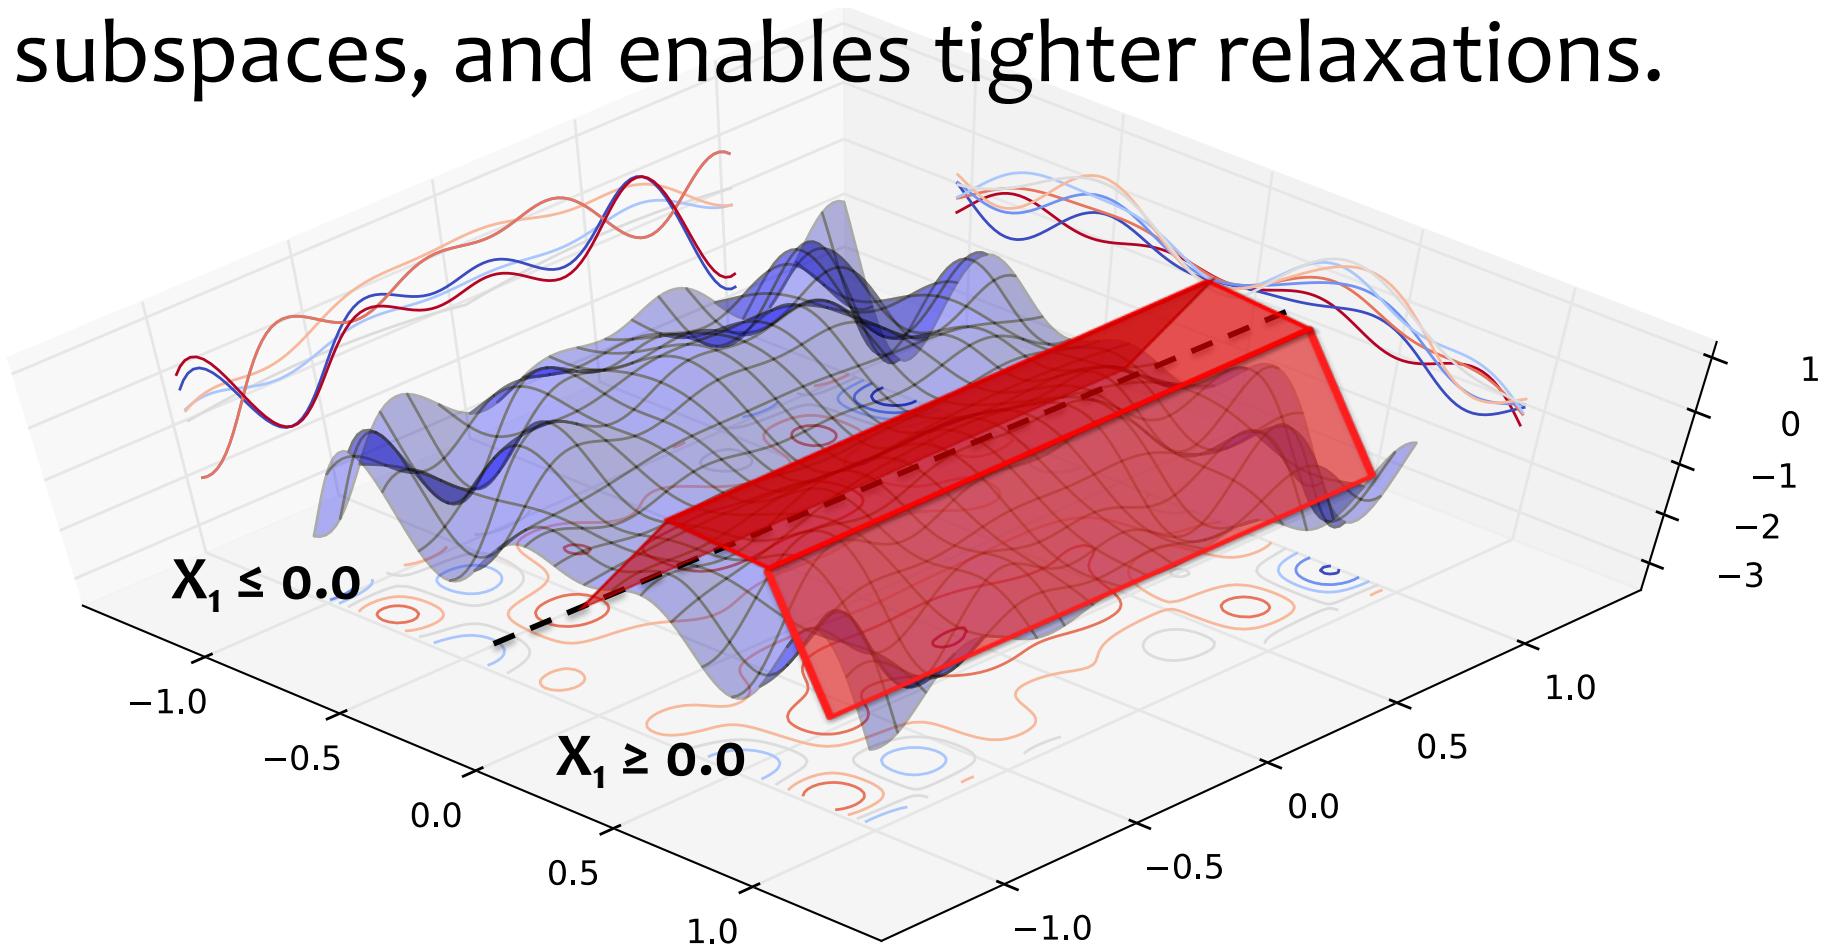
Background: Nonconvex Global Optimization

Branching: partitions the search space into subspaces, and enables tighter relaxations.



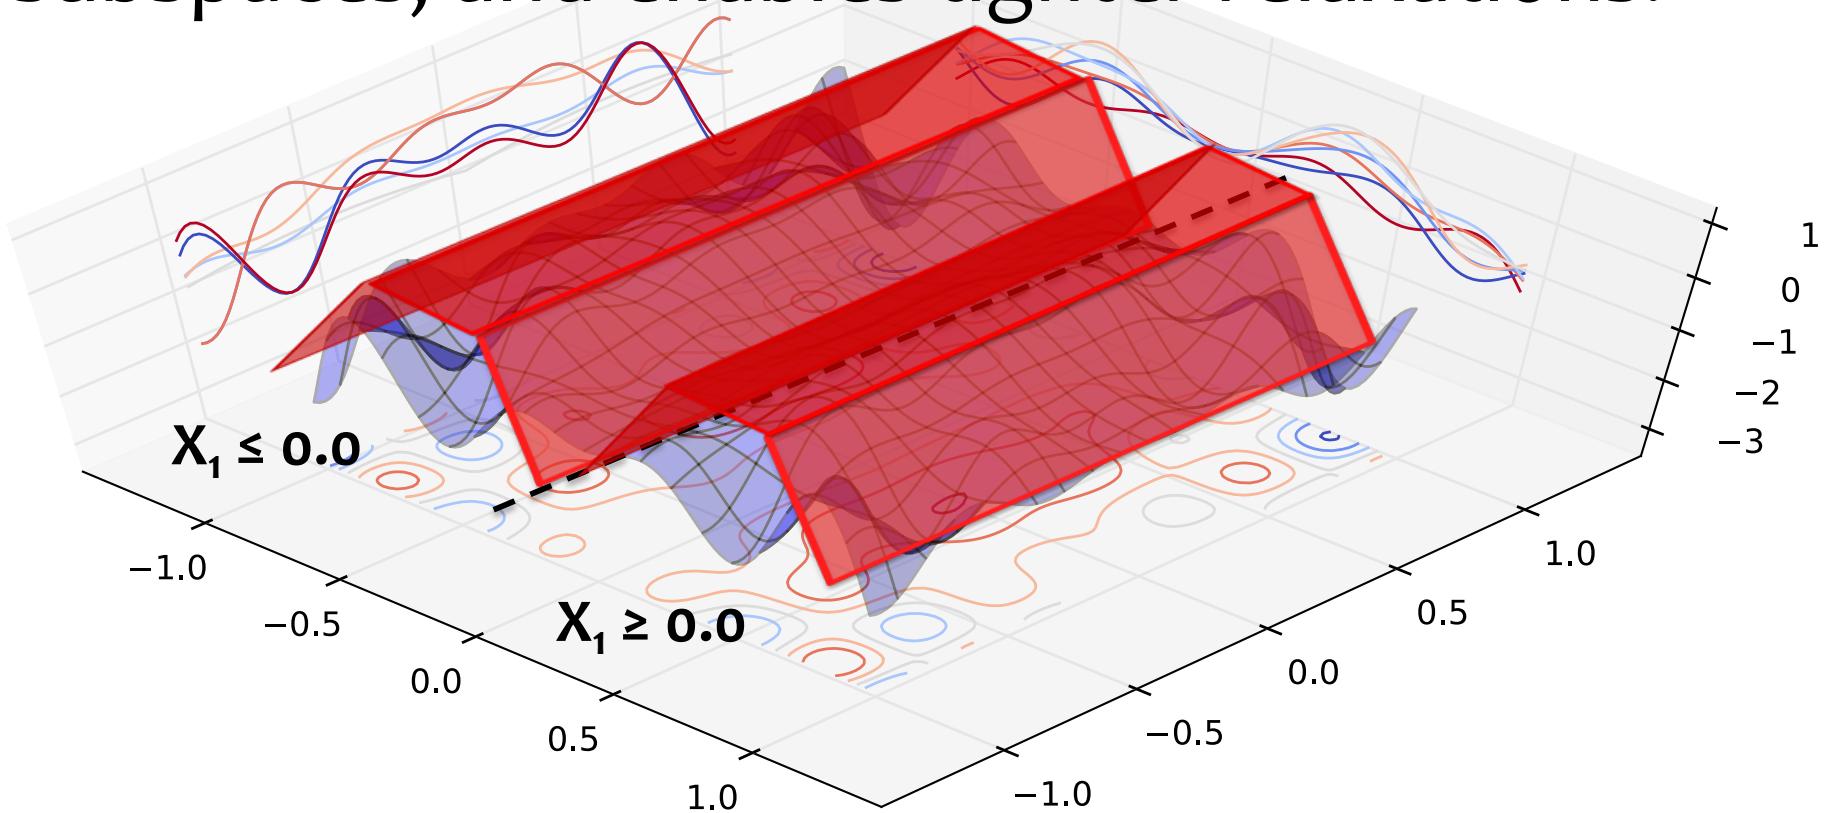
Background: Nonconvex Global Optimization

Branching: partitions the search space into subspaces, and enables tighter relaxations.



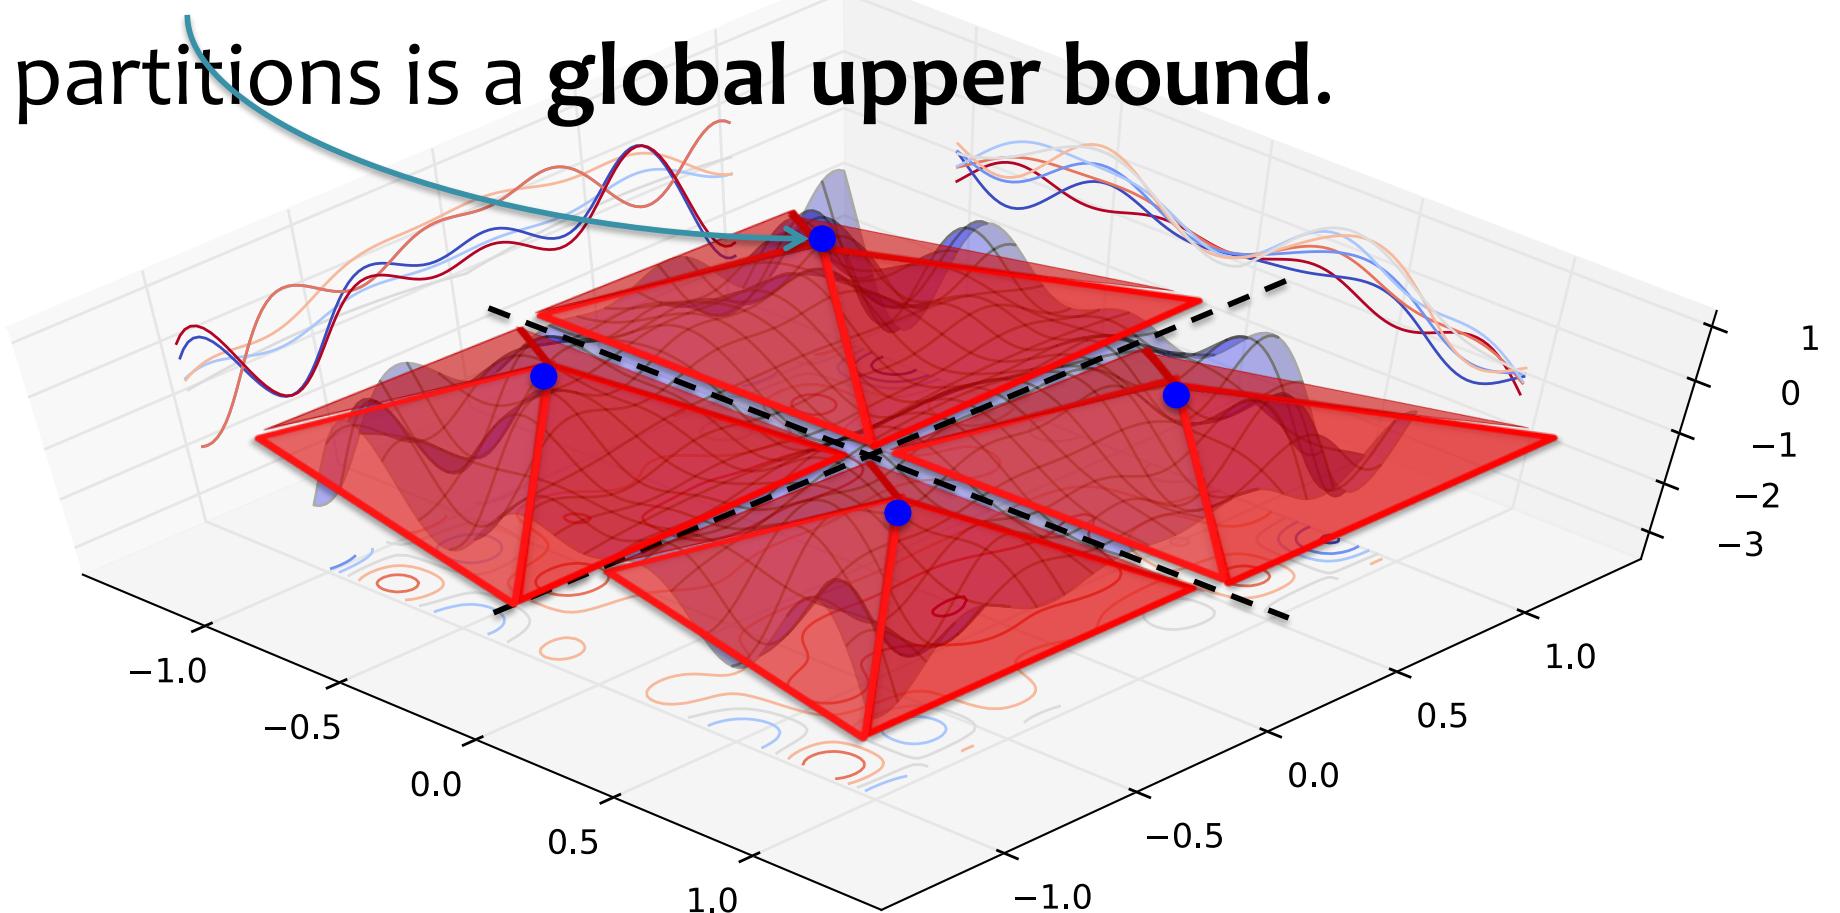
Background: Nonconvex Global Optimization

Branching: partitions the search space into subspaces, and enables tighter relaxations.



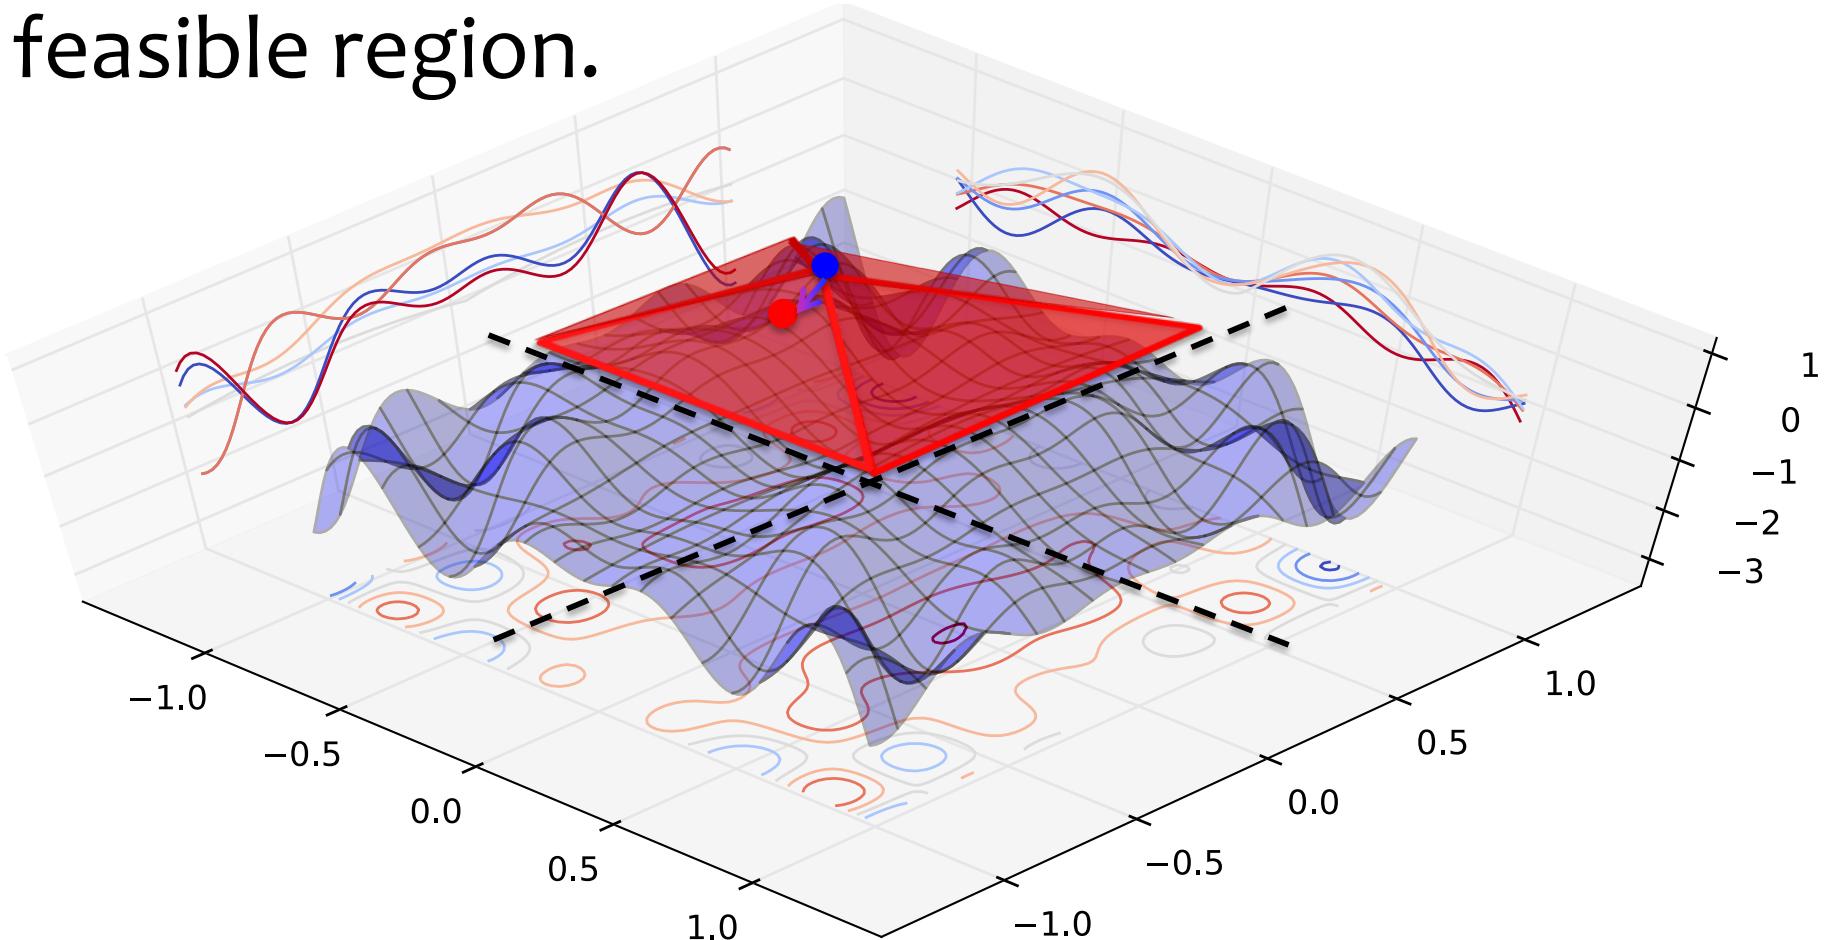
Background: Nonconvex Global Optimization

The **max** of all relaxed solutions for each of the partitions is a **global upper bound**.



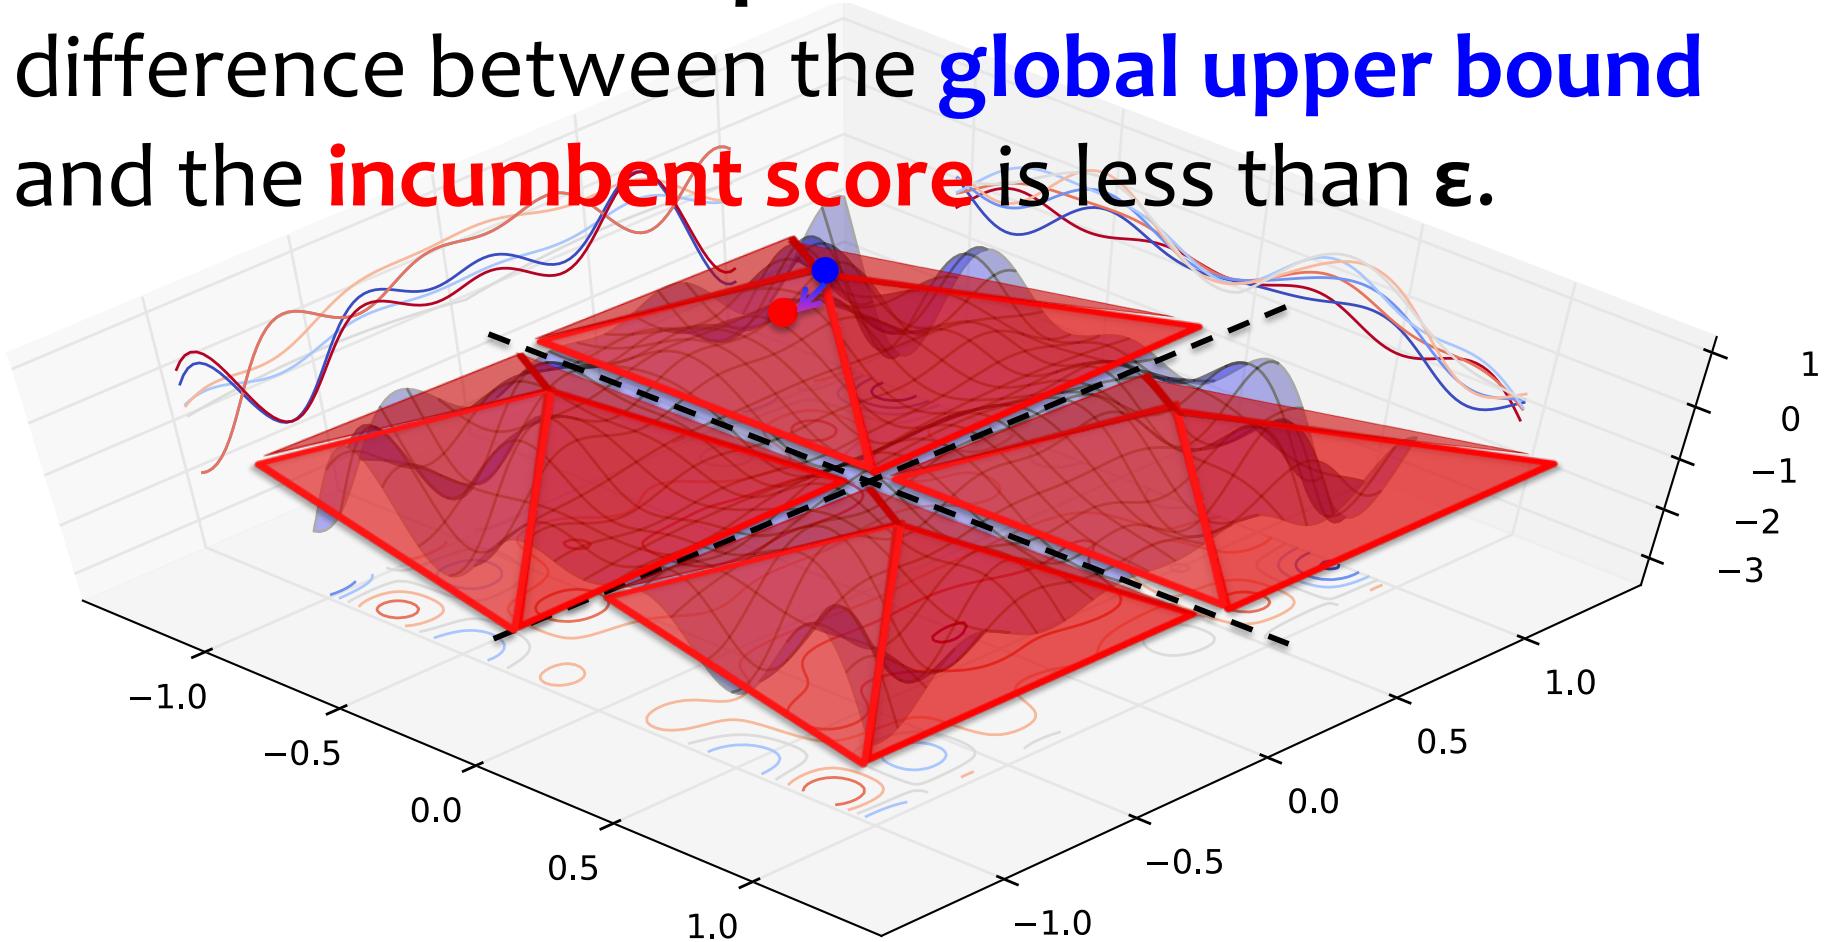
Background: Nonconvex Global Optimization

We can **project** a relaxed solution onto the feasible region.



Background: Nonconvex Global Optimization

The **incumbent** is ϵ -optimal if the relative difference between the **global upper bound** and the **incumbent score** is less than ϵ .



How much should we subdivide?

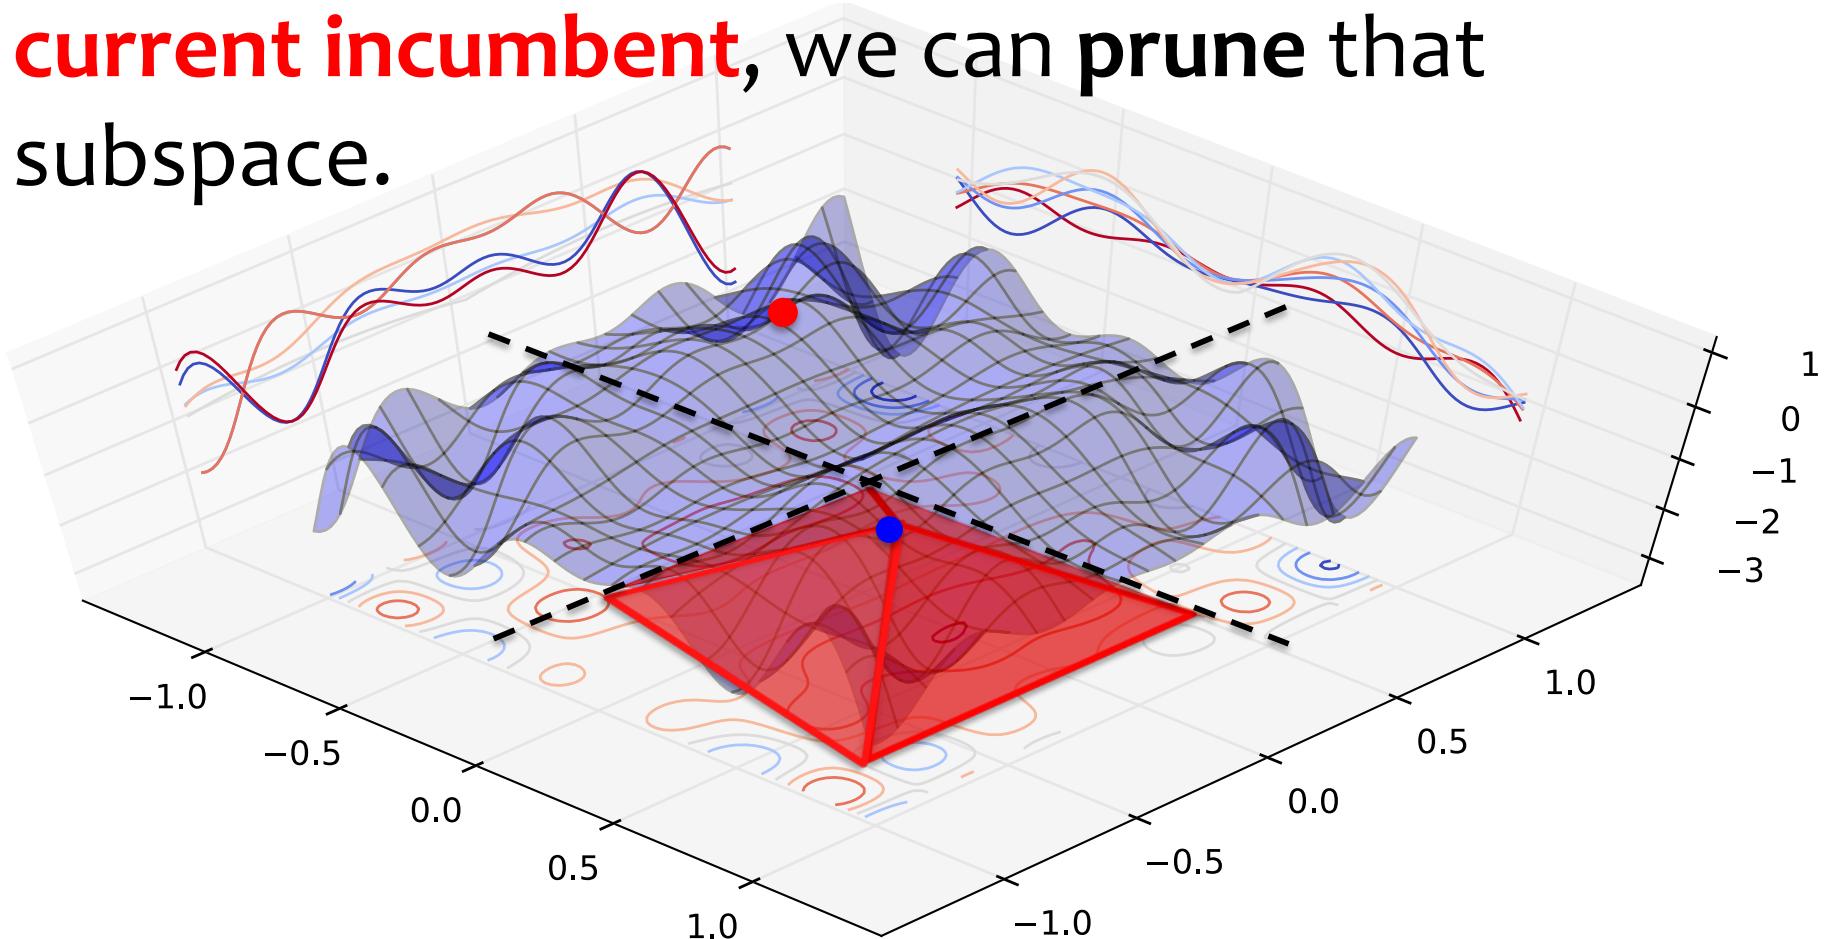
How much should we subdivide?

BRANCH-AND-BOUND

- Method for **recursively subdividing** the search space
- **Subspace order** can be determined heuristically (e.g. best-first search with depth-first plunging)
- **Prunes** subspaces that can't yield better solutions

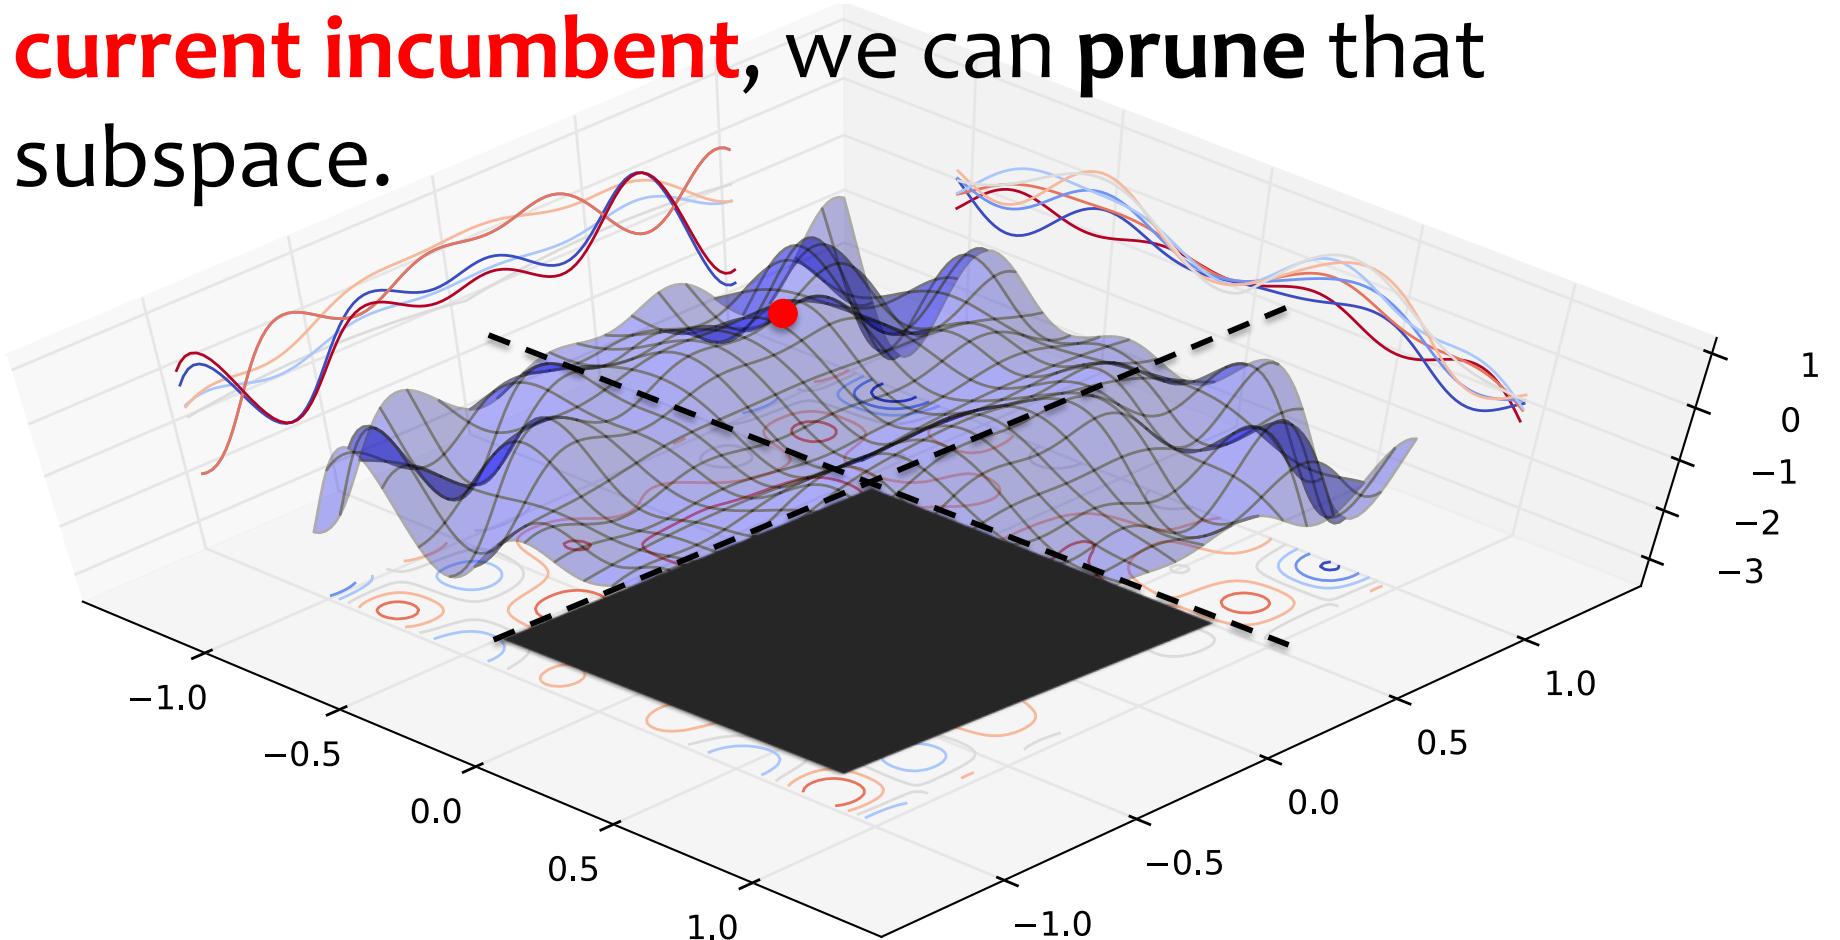
Background: Nonconvex Global Optimization

If the **subspace upper bound** is worse than the **current incumbent**, we can **prune** that subspace.



Background: Nonconvex Global Optimization

If the **subspace upper bound** is worse than the **current incumbent**, we can **prune** that subspace.



Limitations:

Branch-and-Bound

Branch-and-bound

- Kind of tricky to get it right...
 - Lots of hyperparameters to tune
- Curse of dimensionality kicks in quickly
- Applications to problems other than ILP (e.g. QP) rather unsuccessful
 - Nonconvex quadratic optimization by LP-based branch-and-bound usually fails with more than 80 variables (Burer and Vandenbussche, 2009)

BRANCH-AND-BOUND INGREDIENTS

Mathematical Program

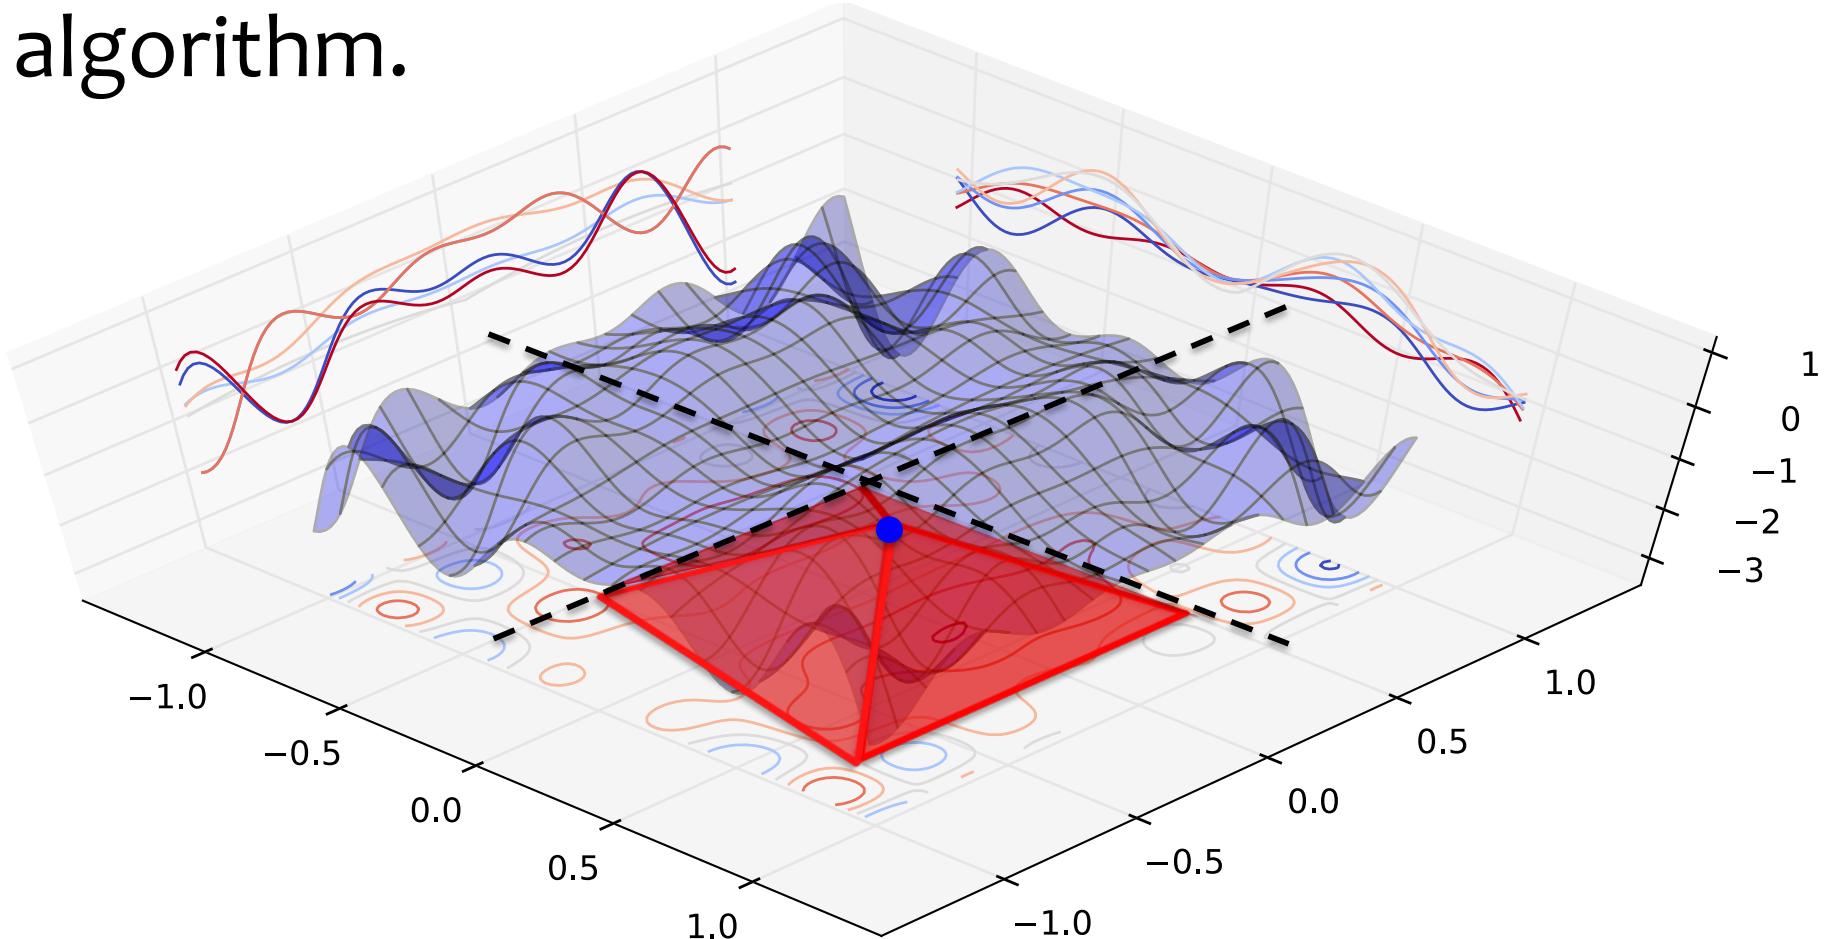
Relaxation

Projection

(Branch-and-Bound Search Heuristics)

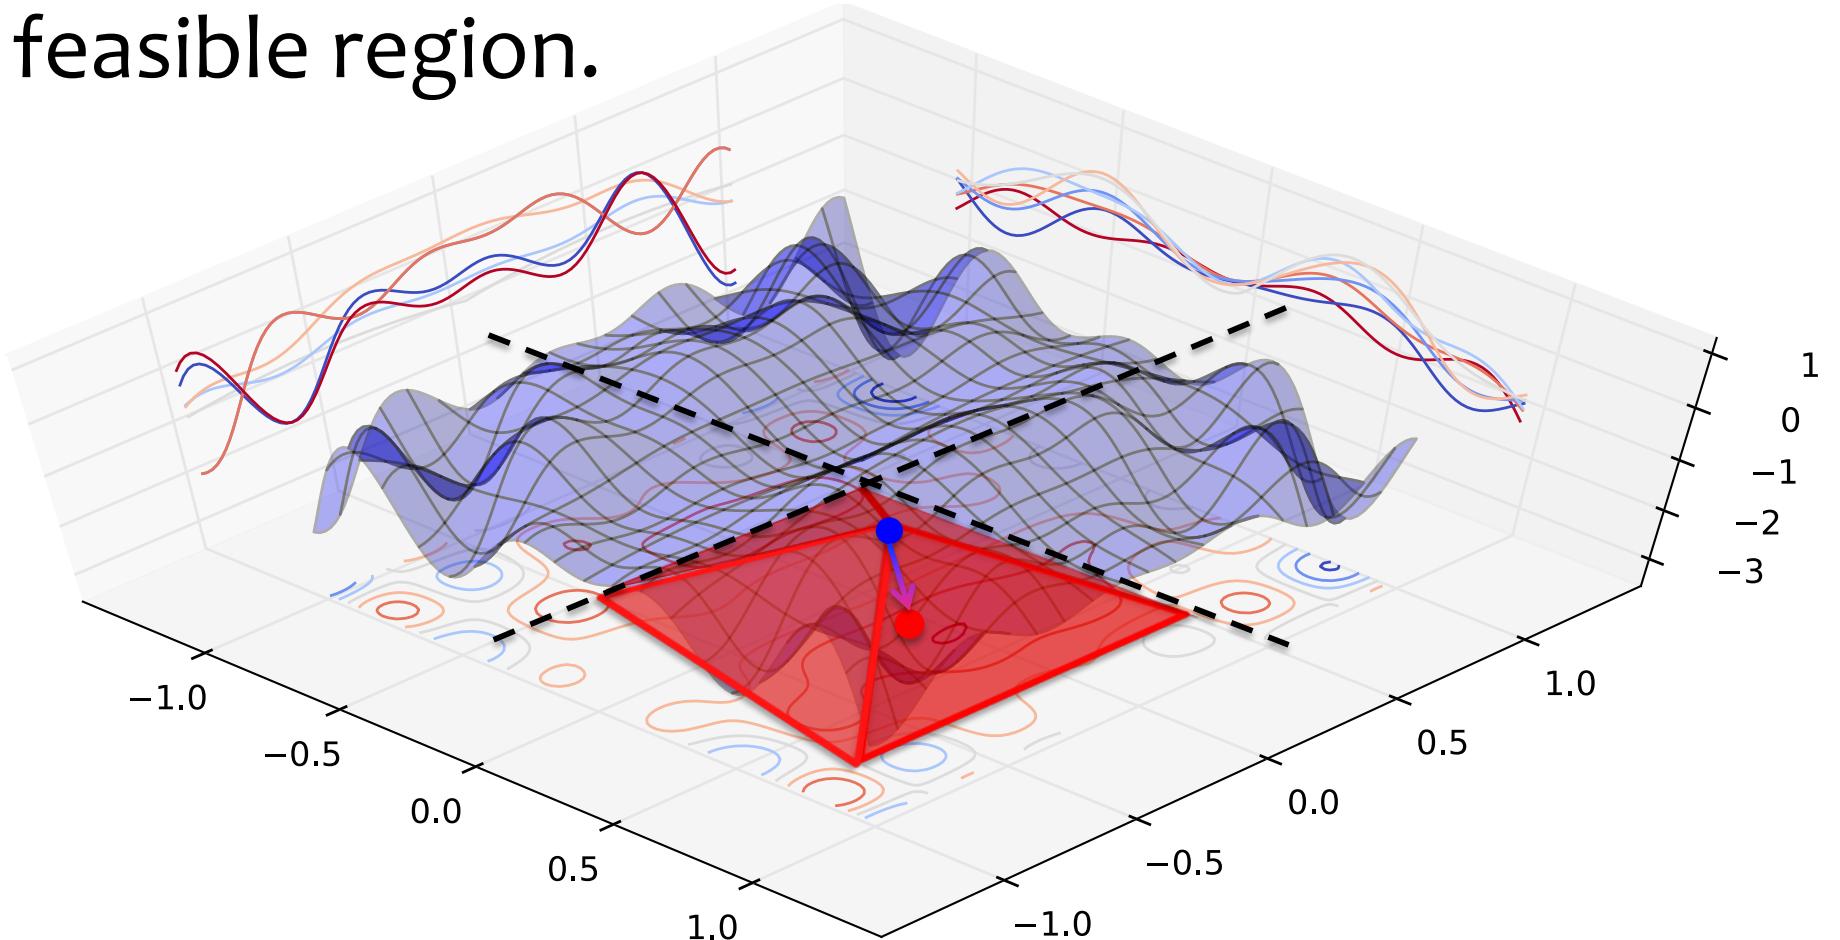
Background: Nonconvex Global Optimization

We solve the **relaxation** using the Simplex algorithm.



Background: Nonconvex Global Optimization

We can **project** a relaxed solution onto the feasible region.



Integer Linear Programming

Whiteboard

- Branch and bound for an ILP in 2D

Branch and Bound

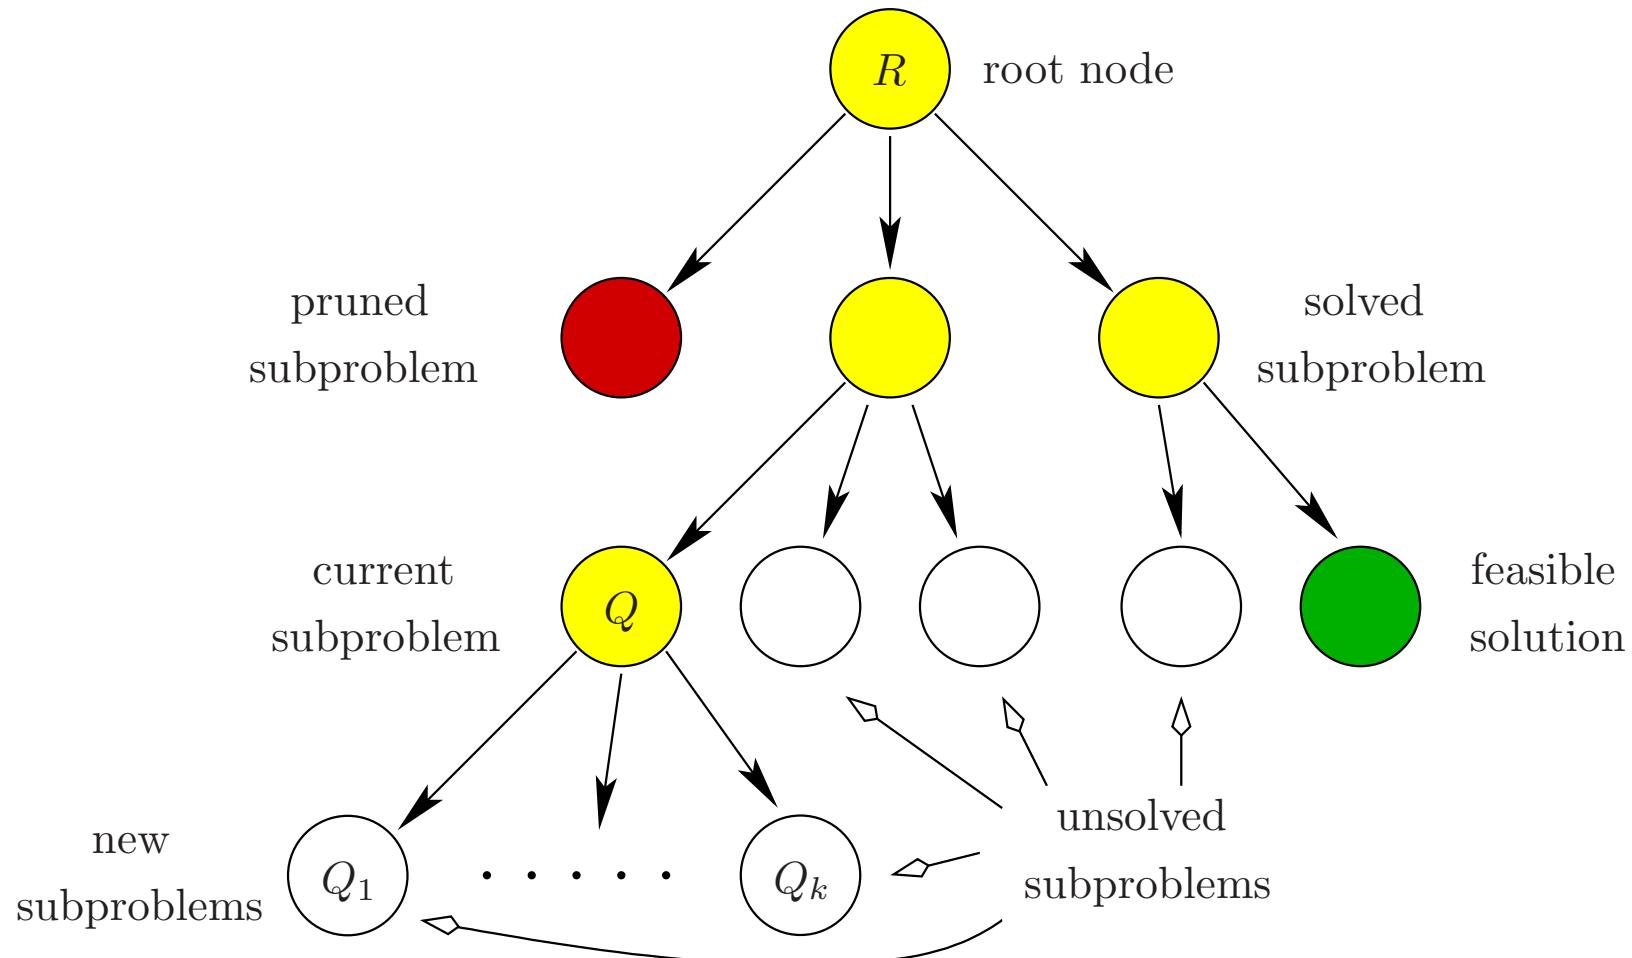
Algorithm 2.1 Branch-and-bound

Input: Minimization problem instance R .

Output: Optimal solution x^* with value c^* , or conclusion that R has no solution, indicated by $c^* = \infty$.

1. Initialize $\mathcal{L} := \{R\}$, $\hat{c} := \infty$. [init]
 2. If $\mathcal{L} = \emptyset$, stop and return $x^* = \hat{x}$ and $c^* = \hat{c}$. [abort]
 3. Choose $Q \in \mathcal{L}$, and set $\mathcal{L} := \mathcal{L} \setminus \{Q\}$. [select]
 4. Solve a relaxation Q_{relax} of Q . If Q_{relax} is empty, set $\check{c} := \infty$. Otherwise, let \check{x} be an optimal solution of Q_{relax} and \check{c} its objective value. [solve]
 5. If $\check{c} \geq \hat{c}$, goto Step 2. [bound]
 6. If \check{x} is feasible for R , set $\hat{x} := \check{x}$, $\hat{c} := \check{c}$, and goto Step 2. [check]
 7. Split Q into subproblems $Q = Q_1 \cup \dots \cup Q_k$, set $\mathcal{L} := \mathcal{L} \cup \{Q_1, \dots, Q_k\}$, and goto Step 2. [branch]
-

Branch and Bound



Branch and Bound

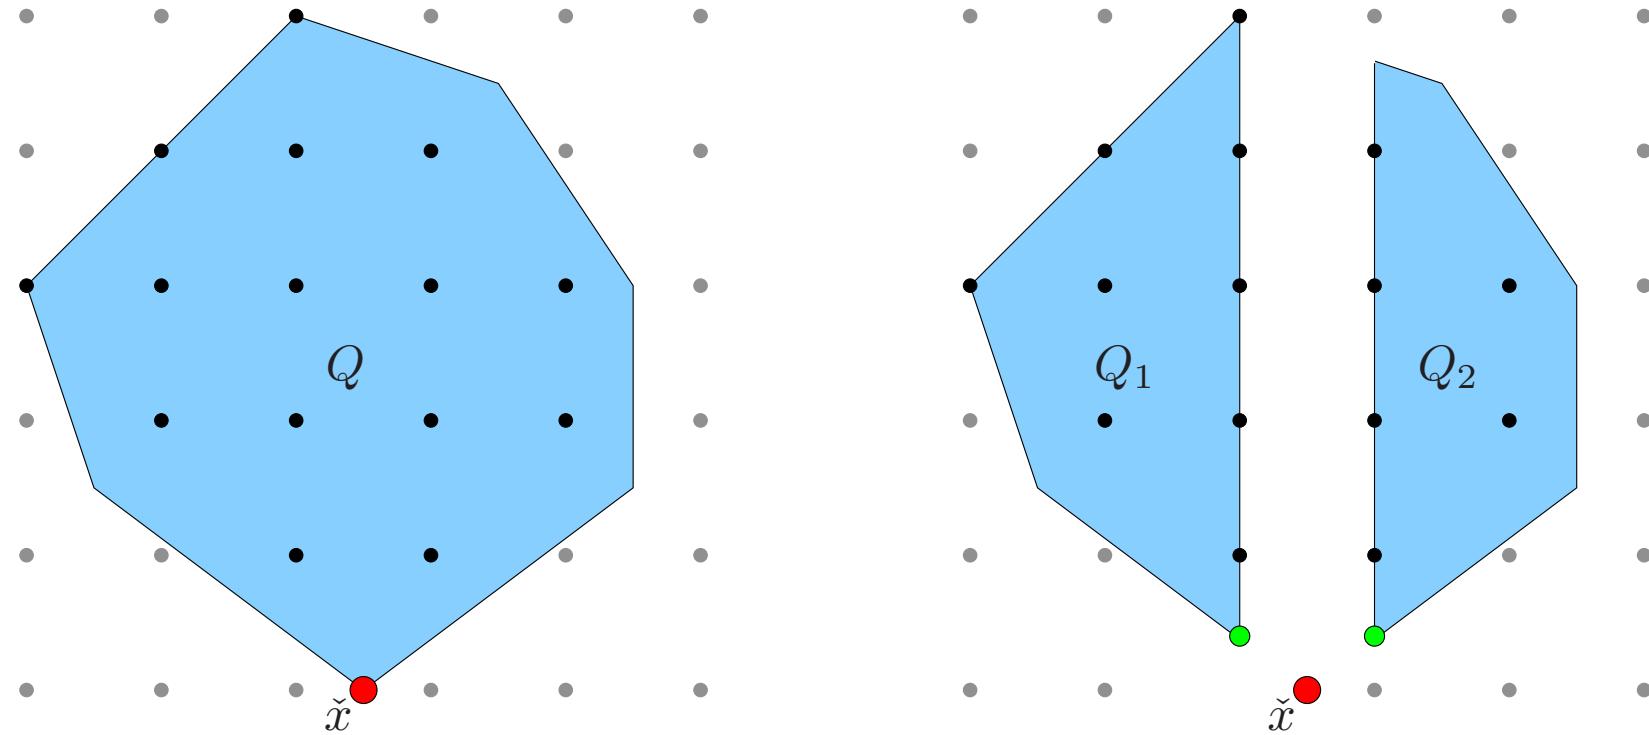


Figure 2.2. LP based branching on a single fractional variable.

Case #1: Binary Variables

MAP INFERENCE AS MATHEMATICAL PROGRAMMING

Recall...

Exact Inference

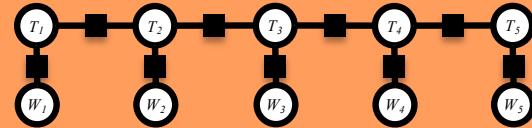
1. Data

$$\mathcal{D} = \{\boldsymbol{x}^{(n)}\}_{n=1}^N$$



2. Model

$$p(\boldsymbol{x} \mid \boldsymbol{\theta}) = \frac{1}{Z(\boldsymbol{\theta})} \prod_{C \in \mathcal{C}} \psi_C(\boldsymbol{x}_C)$$



3. Objective

$$\ell(\boldsymbol{\theta}; \mathcal{D}) = \sum_{n=1}^N \log p(\boldsymbol{x}^{(n)} \mid \boldsymbol{\theta})$$

5. Inference

1. Marginal Inference

$$p(x_C) = \sum_{\boldsymbol{x}' : \boldsymbol{x}'_C = \boldsymbol{x}_C} p(\boldsymbol{x}' \mid \boldsymbol{\theta})$$

2. Partition Function

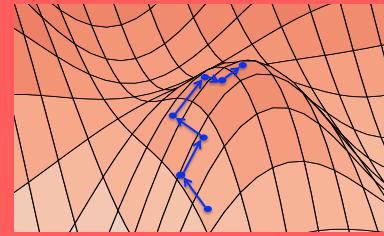
$$Z(\boldsymbol{\theta}) = \sum_{\boldsymbol{x}} \prod_{C \in \mathcal{C}} \psi_C(\boldsymbol{x}_C)$$

3. MAP Inference

$$\hat{\boldsymbol{x}} = \operatorname{argmax}_{\boldsymbol{x}} p(\boldsymbol{x} \mid \boldsymbol{\theta})$$

4. Learning

$$\boldsymbol{\theta}^* = \operatorname{argmax}_{\boldsymbol{\theta}} \ell(\boldsymbol{\theta}; \mathcal{D})$$



5. Inference

Three Tasks:

1. Marginal Inference

Compute marginals of variables and cliques

$$p(x_i) = \sum_{\mathbf{x}': x'_i = x_i} p(\mathbf{x}' | \boldsymbol{\theta}) \quad \mid \quad p(\mathbf{x}_C) = \sum_{\mathbf{x}': \mathbf{x}'_C = \mathbf{x}_C} p(\mathbf{x}' | \boldsymbol{\theta})$$

2. Partition Function

Compute the normalization constant

$$Z(\boldsymbol{\theta}) = \sum_{\mathbf{x}} \prod_{C \in \mathcal{C}} \psi_C(\mathbf{x}_C)$$

3. MAP Inference (NP-Hard in the general case)

Compute variable assignment with highest probability

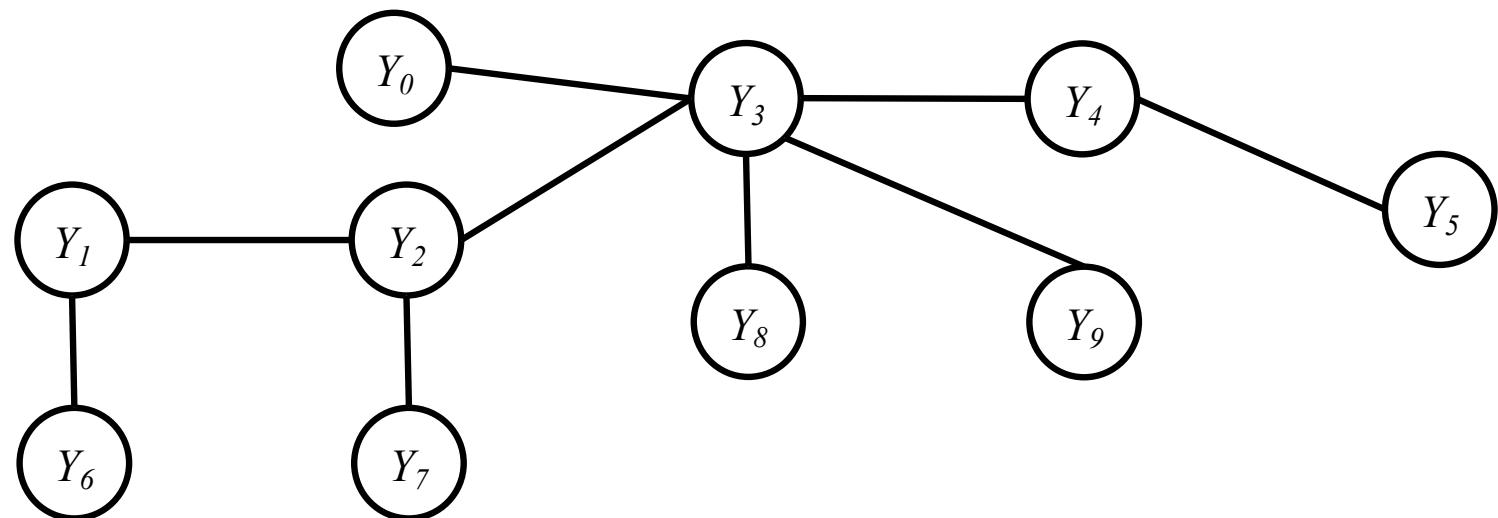
$$\hat{\mathbf{x}} = \operatorname{argmax}_{\mathbf{x}} p(\mathbf{x} | \boldsymbol{\theta})$$

MAP Inference

Suppose we want to predict the highest likelihood structure y , given observations x and parameters w .

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y}} \log p_w(\mathbf{y}|x)$$

$$= \operatorname{argmax}_{\mathbf{y}} \sum_j \mathbf{w}^T f_{\text{node}}(x_j, y_j) + \sum_{j,k} \mathbf{w}^T f_{\text{edge}}(\mathbf{x}_{jk}, y_j, y_k)$$



MAP Inference

Suppose we want to predict the highest likelihood structure y , given observations x and parameters w .

$$\begin{aligned}\hat{\mathbf{y}} &= \operatorname{argmax}_{\mathbf{y}} \log p_w(\mathbf{y}|\mathbf{x}) \\ &= \operatorname{argmax}_{\mathbf{y}} \sum_j \mathbf{w}^T f_{\text{node}}(x_j, y_j) + \sum_{j,k} \mathbf{w}^T f_{\text{edge}}(\mathbf{x}_{jk}, y_j, y_k)\end{aligned}$$

Idea:

1. Reformulate the problem as an integer linear program (ILP) – note that this is just going to be a new way of writing down the problem: $\mathbf{y} \rightarrow \mathbf{z}$
2. Then remove the integer constraints (i.e. solve the linear program (LP) relaxation)

Lemma: (Wainwright et al., 2002) If there is a unique MAP assignment, the LP relaxation of the ILP above is guaranteed to have an integer solution, which is exactly the MAP solution!

Integer Linear Programming

Whiteboard

- MAP Inference for a Binary Pairwise MRF as an ILP