



# 10-708 Probabilistic Graphical Models

Machine Learning Department  
School of Computer Science  
Carnegie Mellon University



## Variational Autoencoders

Matt Gormley  
Lecture 20  
April 12, 2021

# Reminders

- **Quiz 2**
  - Wed, Apr 14, during lecture time
- **HW5 Recitation**
  - Wed, Apr. 14 at 7pm
- **Homework 5: Variational Inference**
  - Out: Thu, Apr. 8
  - Due: Wed, Apr. 21 at 11:59pm
- **Project Midway Milestones:**
  - **Midway Poster Session:**  
Tue, Apr. 27 at 6:30pm – 8:30pm
  - **Midway Executive Summary**  
Due: Tue, Apr. 27 at 11:59pm
  - **New requirement: must have baseline results**

# **QUIZ 2 LOGISTICS**

# Quiz 2

- **Time / Location**
  - **Time:** In-Class Quiz  
**Wed, Apr. 14 during lecture time**
  - **Location:** The same Zoom meeting as lecture/recitation.  
Please arrive online early.
  - Please watch Piazza carefully for announcements.
- **Logistics**
  - Covered material: Lecture 9 – Lecture 15  
(and unavoidably some material from Lectures 1 – 8)
  - Format of questions:
    - Multiple choice
    - True / False (with justification)
    - Derivations
    - Short answers
    - Interpreting figures
    - Implementing algorithms on paper
    - Drawing
  - No electronic devices
  - You are allowed to **bring** one 8½ x 11 sheet of notes (front and back)

# Quiz 2

- **Advice (for before the exam)**
  - Try out the Gradescope quiz-style interface in the “Fake Quiz” now available
- **Advice (for during the exam)**
  - Solve the easy problems first (e.g. multiple choice before derivations)
    - if a problem seems extremely complicated you’re likely missing something
  - Don’t leave any answer blank!
  - If you make an assumption, write it down
  - If you look at a question and don’t know the answer:
    - we probably haven’t told you the answer
    - but we’ve told you enough to work it out
    - imagine arguing for some answer and see if you like it

# Topics for Quiz 1

- Graphical Model Representation
  - Directed GMs vs. Undirected GMs vs. Factor Graphs
  - Bayesian Networks vs. Markov Random Fields vs. Conditional Random Fields
- Graphical Model Learning
  - Fully observed Bayesian Network learning
  - Fully observed MRF learning
  - Fully observed CRF learning
  - Parameterization of a GM
  - Neural potential functions
- Exact Inference
  - Three inference problems:
    - (1) marginals
    - (2) partition function
    - (3) most probably assignment
  - Variable Elimination
  - Belief Propagation (sum-product and max-product)

# Topics for Quiz 2

- Learning for Structure Prediction
  - Structured Perceptron
  - Structured SVM
  - Neural network potentials
- (Approximate) MAP Inference
  - MAP Inference via MILP
  - MAP Inference via LP relaxation
- Approximate Inference by Sampling
  - Monte Carlo Methods
  - Gibbs Sampling
  - Metropolis-Hastings
  - Markov Chains and MCMC
- Parameter Estimation
  - Bayesian inference
  - Topic Modeling

Q&A



# **AUTOENCODERS**

# Unsupervised Pre-training

- **Idea: (Two Steps)**

- Use supervised learning, but **pick a better starting point**
- **Train each level** of the model in a **greedy** way

1. Unsupervised Pre-training

- Use **unlabeled** data
- Work bottom-up
  - Train hidden layer 1. Then fix its parameters.
  - Train hidden layer 2. Then fix its parameters.
  - ...
  - Train hidden layer n. Then fix its parameters.


2. Supervised Fine-tuning

- Use **labeled** data to train following “Idea #1”
- Refine the features by backpropagation so that they become tuned to the end-task

# Unsupervised Pre-training

## Unsupervised pre-training of the first layer:

- What should it predict?
- What else do we observe?
- **The input!**




# Auto-Encoders

**Unsupervised pre-training of the first layer:**

- What should it predict?
- What else do we observe?
- **The input!**

**This topology defines an Auto-encoder.**




# Auto-Encoders

Key idea: Encourage  $z$  to give small reconstruction error:

- $x'$  is the *reconstruction* of  $x$
- $\text{Loss} = ||x - \text{DECODER}(\text{ENCODER}(x))||^2$
- Train with the same backpropagation algorithm for 2-layer Neural Networks with  $x_m$  as both input and output.

DECODER:  $x' = h(W'z)$


ENCODER:  $z = h(Wx)$



# Unsupervised Pre-training

## Unsupervised pre-training

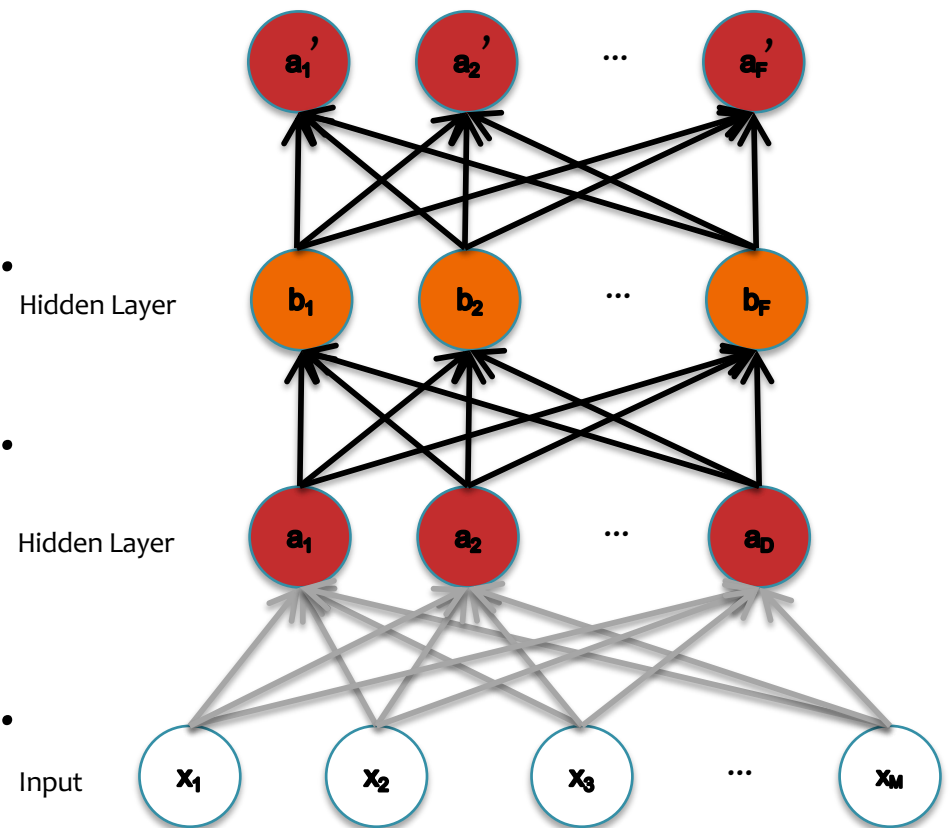
- Work bottom-up
  - Train hidden layer 1.  
Then fix its parameters.
  - Train hidden layer 2.  
Then fix its parameters.
  - ...
  - Train hidden layer  $n$ .  
Then fix its parameters.



# Unsupervised Pre-training

## Unsupervised pre-training


- Work bottom-up
  - Train hidden layer 1. Then fix its parameters.
  - Train hidden layer 2. Then fix its parameters.
  - ...
  - Train hidden layer  $n$ . Then fix its parameters.



# Unsupervised Pre-training

## Unsupervised pre-training

- Work bottom-up
  - Train hidden layer 1. Then fix its parameters.
  - Train hidden layer 2. Then fix its parameters.
  - ...
  - Train hidden layer  $n$ . Then fix its parameters.





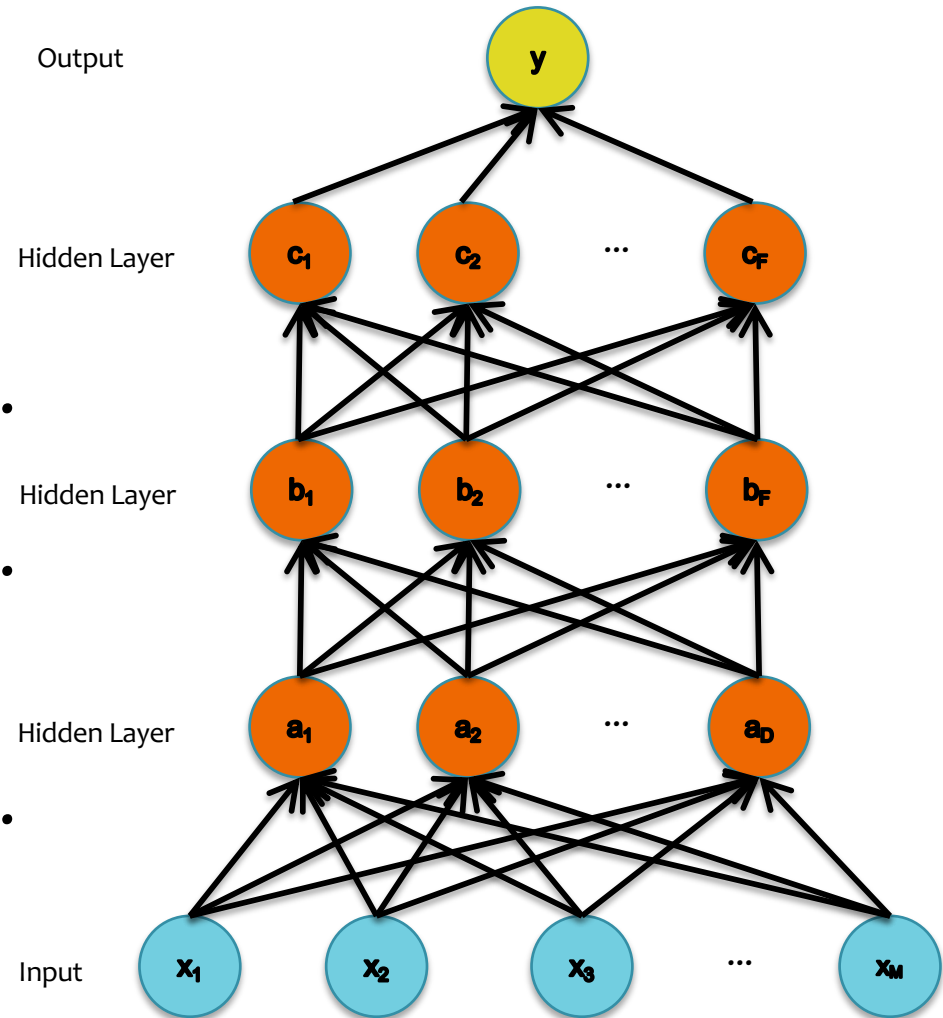
# Unsupervised Pre-training

## Unsupervised pre-training

- Work bottom-up
  - Train hidden layer 1. Then fix its parameters.
  - Train hidden layer 2. Then fix its parameters.
  - ...
  - Train hidden layer n. Then fix its parameters.

## Supervised fine-tuning

Backprop and update all parameters



# Deep Network Training

- **Idea #1:**

1. Supervised fine-tuning only

- **Idea #2:**


1. Supervised layer-wise pre-training
2. Supervised fine-tuning

- **Idea #3:**

1. Unsupervised layer-wise pre-training
2. Supervised fine-tuning


# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)



# Comparison on MNIST

- Results from Bengio et al. (2006) on MNIST digit classification task
- Percent error (lower is better)



# VARIATIONAL AUTOENCODERS

# Why VAEs?

- **Autoencoders:**


- learn a low dimensional representation of the input, but hard to work with as a generative model
- one of the key limitations of autoencoders is that we have no way of **sampling** from them!

- **Variational autoencoders (VAEs)**

- by contrast learn a continuous latent space that is **easy to sample from!**
- can **generate** new data (e.g. images) by sampling from the learned generative model

# Variational Autoencoders



$$p_{\phi}(\mathbf{x}, \mathbf{z})$$



$$\mathbf{z} \sim \text{Gaussian}(\mathbf{0}, \mathbf{I})$$

## Graphical Model Perspective



- The DGM diagram shows that the VAE model is quite simple as a graphical model (ignoring the neural net details that give rise to  $\mathbf{x}$ )
- Sampling from the model is easy:
  - Consider a DGM where  $\mathbf{x} = g_{\phi}(\mathbf{z}/10 + \mathbf{z}/\|\mathbf{z}\|)$  (i.e. we don't use parameters  $\phi$ )
  - Then we can draw samples of  $\mathbf{z}$  and directly convert them to values  $\mathbf{x}$
- **Key idea of VAE:** define  $g_{\phi}(\mathbf{z})$  as a neural net and learn  $\phi$  from data



# Variational Autoencoders

## Neural Network Perspective

- We can view a variational autoencoder (VAE) as an autoencoder consisting of two neural networks
- VAEs (as encoders) define two distributions:
  - **encoder:**  $q_{\theta}(z | x)$
  - **decoder:**  $p_{\phi}(x | z)$
- Parameters  $\theta$  and  $\phi$  are neural network parameters (i.e.  $\theta$  are not the variational parameters)







# Variational Autoencoders

## Graphical Model Perspective

- We can also view the VAE from the perspective of variational inference
- In this case we have two distributions:
  - **model:**  $p_{\phi}(z | x)$
  - **variational approximation:**  $q_{\lambda=f(x; \theta)}(z | x)$
- We have the same model parameters  $\phi$
- The variational parameters  $\lambda$  are a function of NN parameters  $\theta$



# Variational Autoencoders

## ***Whiteboard***

- Variational Autoencoder = VAE
- VAE as a Probability Model
- Parameterizing the VAE with Neural Nets
- Variational EM for VAEs

# Reparameterization Trick

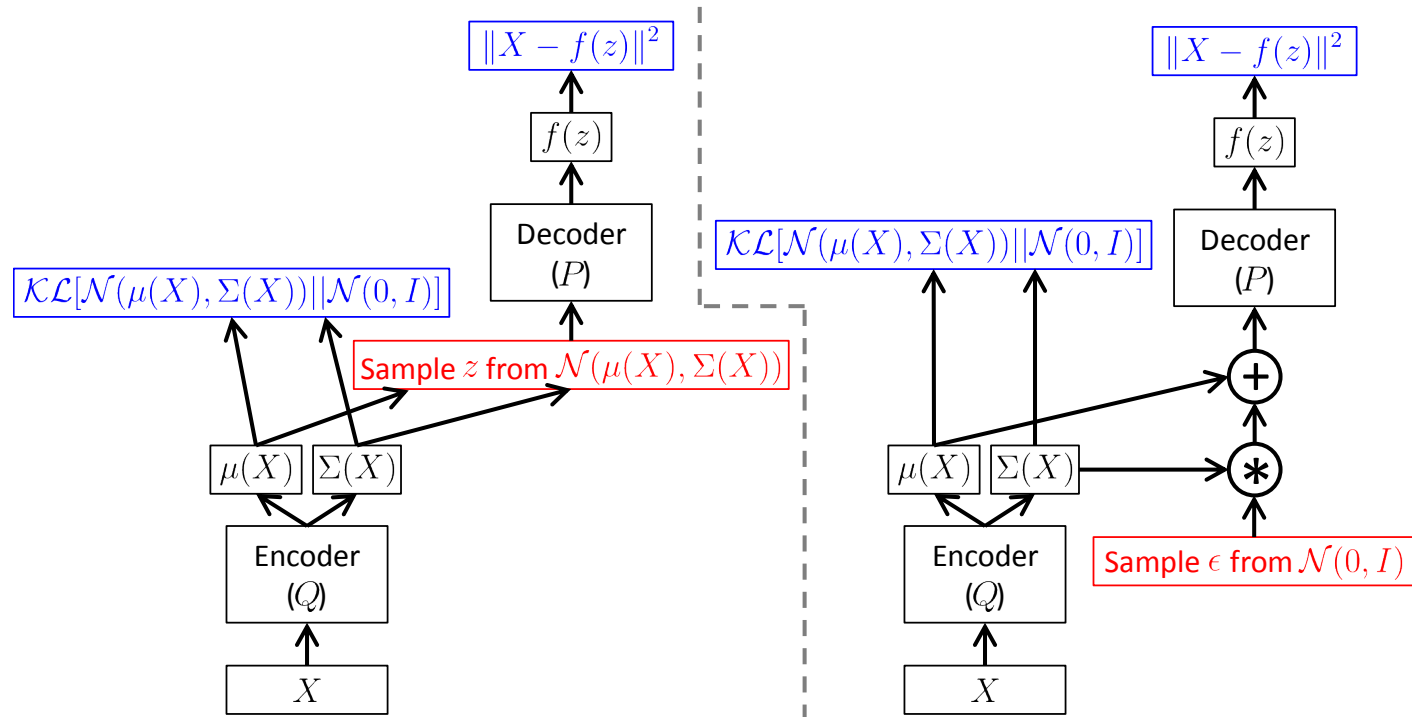


Figure 4: A training-time variational autoencoder implemented as a feed-forward neural network, where  $P(X|z)$  is Gaussian. Left is without the “reparameterization trick”, and right is with it. Red shows sampling operations that are non-differentiable. Blue shows loss layers. The feedforward behavior of these networks is identical, but backpropagation can be applied only to the right network.