

22AIE304 Deep Learning Lab Sheet 2

Fifth Semester BTech CSE(AI)

Department of Computer Science and Engineering

Amrita School of Computing

Neural Networks with McCulloch-Pitts Neurons and Perceptron

Exercise 1: Use MP neurons to build a simple neural network that performs logical operations OR, AND and NOR

```
import numpy as np

def MP_neuron(weights, bias, inputs):
    z = np.dot(weights, inputs) + bias
    return 1 if z >= 0 else 0

inputs = [[0, 0], [0, 1], [1, 0], [1, 1]]

# OR
weights = np.array([1, 1])
bias = -0.5

print("OR Operation")
for inp in inputs:
    print(f"Input: {inp}, Output: {MP_neuron(weights, bias, inp)}")

# AND
weights = np.array([1, 1])
bias = -1.5

print("\n\nAND Operation")
for inp in inputs:
    print(f"Input: {inp}, Output: {MP_neuron(weights, bias, inp)}")

# NOR
weights = np.array([-1, -1])
bias = 0.5

print("\n\nNOR Operation")
for inp in inputs:
    print(f"Input: {inp}, Output: {MP_neuron(weights, bias, inp)}")
```

➡ OR Operation

```
Input: [0, 0], Output: 0
Input: [0, 1], Output: 1
Input: [1, 0], Output: 1
Input: [1, 1], Output: 1
```

AND Operation

```
Input: [0, 0], Output: 0
Input: [0, 1], Output: 0
Input: [1, 0], Output: 0
Input: [1, 1], Output: 1
```

NOR Operation

```
Input: [0, 0], Output: 1
Input: [0, 1], Output: 0
```

Input: [1, 0], Output: 0
 Input: [1, 1], Output: 0

✓ **NOR cannot be learned by McCulloch-Pitts (MP) neuron, as it can only learn linearly seperable data, while NOR data points are Non-Linear in nature.**

Exercise 2: Implement an MP neuron for a binary classification problem using a breast cancer dataset.

- Analyze the effects of scaling on MP Neuron's decision-making process and accuracy. Apply different scaling techniques (min-max normalization, standardization) to the breast cancer dataset features. Train the MP Neuron with these scaled features and compare the model's performance with unscaled data.
- Compare the MP Neuron model's performance with a logistic regression model in accuracy.

```
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler
import numpy as np

data = load_breast_cancer()
X, y = data.data, data.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

scaler_minmax = MinMaxScaler()
X_train_minmax = scaler_minmax.fit_transform(X_train)
X_test_minmax = scaler_minmax.transform(X_test)

scaler_standard = StandardScaler()
X_train_std = scaler_standard.fit_transform(X_train)
X_test_std = scaler_standard.transform(X_test)

def MP_neuron(weights, bias, inputs):
    weighted_sum = np.dot(weights, inputs) + bias
    return 1 if weighted_sum >= 0 else 0

def train_MP_neuron(X_train, y_train, threshold=0):
    weights = np.ones(X_train.shape[1])
    bias = threshold
    predictions = [MP_neuron(weights, bias, x) for x in X_train]
    accuracy = np.mean(predictions == y_train)
    return accuracy

accuracy_unscaled = train_MP_neuron(X_train, y_train)
accuracy_minmax = train_MP_neuron(X_train_minmax, y_train)
accuracy_std = train_MP_neuron(X_train_std, y_train)

print(f"Accuracy with unscaled data: {accuracy_unscaled}")
print(f"Accuracy with min-max scaled data: {accuracy_minmax}")
print(f"Accuracy with standardized data: {accuracy_std}")

↗ Accuracy with unscaled data: 0.6256281407035176
Accuracy with min-max scaled data: 0.6256281407035176
Accuracy with standardized data: 0.11557788944723618

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

logreg = LogisticRegression()
logreg.fit(X_train, y_train)
y_pred_unscaled = logreg.predict(X_test)
accuracy_logreg_unscaled = accuracy_score(y_test, y_pred_unscaled)

logreg_minmax = LogisticRegression()
logreg_minmax.fit(X_train_minmax, y_train)
y_pred_minmax = logreg_minmax.predict(X_test_minmax)
accuracy_logreg_minmax = accuracy_score(y_test, y_pred_minmax)
```

```

logreg_std = LogisticRegression()
logreg_std.fit(X_train_std, y_train)
y_pred_std = logreg_std.predict(X_test_std)
accuracy_logreg_std = accuracy_score(y_test, y_pred_std)

print(f"Logistic Regression accuracy with unscaled data: {accuracy_logreg_unscaled}")
print(f"Logistic Regression accuracy with min-max scaled data: {accuracy_logreg_minmax}")
print(f"Logistic Regression accuracy with standardized data: {accuracy_logreg_std}")

```

Logistic Regression accuracy with unscaled data: 0.9707602339181286
 Logistic Regression accuracy with min-max scaled data: 0.9649122807017544
 Logistic Regression accuracy with standardized data: 0.9824561403508771
 /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
 Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
 n_iter_i = _check_optimize_result(

Exercise 3: Implement Perceptron for Breast Cancer Classification

- Load and Explore the Dataset
- Split the Data into Training and Testing Sets
- Standardize the Features : Standardize the features (i.e., transform them to have a mean of 0 and a variance of 1) to ensure faster convergence of the Perceptron algorithm.
- Train the Perceptron Model
- Make Predictions
- Evaluate the Model

```

from sklearn.linear_model import Perceptron
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

data = load_breast_cancer()
X, y = data.data, data.target
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

scaler = StandardScaler()
X_train_std = scaler.fit_transform(X_train)
X_test_std = scaler.transform(X_test)

perceptron = Perceptron()
perceptron.fit(X_train_std, y_train)
y_pred = perceptron.predict(X_test_std)

accuracy = accuracy_score(y_test, y_pred)
print(f"Perceptron Accuracy: {accuracy}")

```

Perceptron Accuracy: 0.9649122807017544

Exercise 4: Implement Perceptron algorithm on the binary Iris dataset and explore its performance by adjusting learning rates and analyzing the weight changes during training

- Understand the Iris Dataset and write a summary of features
- Train/ Test Split
- Implement the Perceptron Algorithm

- d. **Plot Train/Test Accuracy:** Once the model is trained, evaluate the accuracy on both the training and testing datasets. Plot the accuracy for the training and testing data to visualize the model's performance over multiple epochs.
- e. **Experimenting with Learning Rates**
- f. **Run the Perceptron algorithm with different learning rates.**
 - Observe how changing the learning rate impacts the model's ability to converge and its overall accuracy.
 - Interpret the results: Does a higher learning rate lead to faster convergence or instability? Does a lower learning rate affect the speed or quality of the model's learning?
- g. **Visualizing the Weight Changes**
 - During training, the Perceptron's weights are updated in each epoch. To understand how the weights evolve:
 - Create a weight matrix that stores the weight values for each epoch.
 - After each epoch, append the current weights to the matrix.
 - Plot the weights as they change across epochs. This will help visualize how the model adjusts its weights based on the data.
 - Write the Interpretation: After plotting the weight changes, explain how the model's weights stabilize as it learns from the data. Do weights converge?
- h. **(Optional) Create an Animation of Weight Changes**

```

from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Perceptron
from sklearn.metrics import accuracy_score
import numpy as np
import matplotlib.pyplot as plt

iris = load_iris()
X, y = iris.data, iris.target
binary_mask = (y == 0) | (y == 1)
X_binary, y_binary = X[binary_mask], y[binary_mask]

X_train, X_test, y_train, y_test = train_test_split(X_binary, y_binary, test_size=0.3, random_state=42)

scaler = StandardScaler()
X_train_std = scaler.fit_transform(X_train)
X_test_std = scaler.transform(X_test)

learning_rates = [0.01, 0.1, 1]
for lr in learning_rates:
    perceptron = Perceptron(eta0=lr)
    perceptron.fit(X_train_std, y_train)
    y_pred_train = perceptron.predict(X_train_std)
    y_pred_test = perceptron.predict(X_test_std)
    train_acc = accuracy_score(y_train, y_pred_train)
    test_acc = accuracy_score(y_test, y_pred_test)
    print(f"Learning Rate {lr} - Train Accuracy: {train_acc}, Test Accuracy: {test_acc}")

↗ Learning Rate 0.01 - Train Accuracy: 1.0, Test Accuracy: 1.0
  Learning Rate 0.1 - Train Accuracy: 1.0, Test Accuracy: 1.0
  Learning Rate 1 - Train Accuracy: 1.0, Test Accuracy: 1.0

weights = []
for epoch in range(10):
    perceptron.partial_fit(X_train_std, y_train, classes=np.unique(y_train))
    weights.append(perceptron.coef_.copy())

weights = np.array(weights)
for i in range(weights.shape[1]):

```

```
plt.plot(weights[:, i, 0], label=f"Feature {i+1}")  
plt.title("Weight Changes Across Epochs")  
plt.xlabel("Epoch")  
plt.ylabel("Weight Value")  
plt.legend()  
plt.show()
```

